



CS 181 - Artificial Intelligence



Kewei Tu, Xuming He
Fall 2020

Administrative Stuff

▶ Instructors:

	屠可伟	何旭明
Email	tukw@shanghaitech.edu.cn	hexm@shanghaitech.edu.cn
Office	SIST 1A-304B	SIST 1A-304D

- ▶ TA: 蒋承越、贾子夏、乔文汇、陈德弘、吴昊一、张甘霖
 - ▶ Office hours: TBA



Administrative Stuff

▶ Classes

- ▶ Mon/Wed 8:15-9:55am @教学中心301
- ▶ 16 weeks
- ▶ Language: English



▶ Prerequisite

- ▶ Programming Languages, Data Structures and Algorithms
- ▶ Discrete Mathematics
- ▶ Probability and Statistics



Administrative Stuff

- ▶ Main textbook
 - ▶ [AIMA] Artificial Intelligence: A Modern Approach, Stuart Russell and Peter Norvig, 3rd edition, 2009.
 - ▶ [中译版] 人工智能：一种现代的方法（第3版），2013



- ▶ Additional reference books may also be used
-

Administrative Stuff

- ▶ Blackboard
 - ▶ Announcements, homework assignments, slides, etc.
- ▶ Piazza
 - ▶ Discussion and QA
 - ▶ <http://piazza.com/shanghaitech.edu.cn/fall2020/cs181>
- ▶ AutoLab
 - ▶ Programming assignments



Administrative Stuff

- ▶ Grading (tentative)
 - ▶ 6 homework assignments (10%)
 - ▶ 6 programming assignments (25%)
 - ▶ Project (15%): 2nd half of the semester
 - ▶ Midterm exam (25%): in late Oct
 - ▶ Final exam (25%): in late Dec or early Jan



Administrative Stuff

▶ Plagiarism

- ▶ All assignments must be done individually
 - ▶ You may not look at solutions from any other source
 - ▶ You may not share solutions with any other students
 - ▶ Plagiarism detection software will be used on all the programming assignments
- ▶ Way of collaboration
 - ▶ You may discuss together or help another student debug code; however, you cannot dictate or give the exact solution



Administrative Stuff

- ▶ Plagiarism punishment
 - ▶ When one student copies from another student, both students are responsible
 - ▶ Zero point on the assignment or exam in question
 - ▶ Repeated violation will result in an F grade for this course as well as further discipline at the school/university level





A brief overview of AI



What is artificial intelligence?

▶ AI = Robots?



Definition

Artificial Intelligence



Machines (Computers)




Rationality:
Ability to maximize goal
achievement given
available information



Artificial Intelligence

- ▶ AI vs. Human Intelligence
 - ▶ Brains (human minds) are good at rational thinking, but not perfect
 - ▶ “Brains are to intelligence as wings are to flight”



A brief history

- ▶ Lots of early speculation & research
 - ▶ Turing: “Computing Machinery and Intelligence” (1950)

I.—COMPUTING MACHINERY AND INTELLIGENCE

BY A. M. TURING

1. *The Imitation Game.*

I PROPOSE to consider the question, ‘Can machines think?’ This should begin with definitions of the meaning of the terms ‘machine’ and ‘think’. The definitions might be framed so as to reflect so far as possible the normal use of the words, but this attitude is dangerous. If the meaning of the words ‘machine’ and ‘think’ are to be found by examining how they are commonly used it is difficult to escape the conclusion that the meaning and the answer to the question, ‘Can machines think?’ is to be sought in a statistical survey such as a Gallup poll. But this is absurd. Instead of attempting such a definition I shall replace the question by another, which is closely related to it and is expressed




Alan Turing



A brief history

- ▶ Birth (1956)
 - ▶ Dartmouth workshop


Dartmouth Conference: The Founding Fathers of AI




John McCarthy



Marvin Minsky



Claude Shannon



Ray Solomonoff

Alan Newell



Herbert Simon



Arthur Samuel



And several
other people...

A brief history

- ▶ Great expectations (1950s-1960s)
 - ▶ A variety of methodology
 - ▶ e.g., symbolism, connectionism
- ▶ AI winter (1970s)
 - ▶ Downfall of perceptron
 - ▶ Lighthill report
- ▶ Boom (1980s)
 - ▶ Revival of neural networks
- ▶ More scientific methods (1990s-2000s)
 - ▶ Statistical approaches

Very rough timeline




A brief history

- ▶ The past ten years
 - ▶ Rise of big data and deep learning
 - ▶ AI becomes one of the hottest areas in CS
 - ▶ Great interest from industry and public
 - ▶ Many real-world applications




Growth of annually published papers, by topic (1996-2017)

Source: Scopus




Attendance at large conferences (1984-2019)

Source: Conference provided data.




AI startups (U.S., Jan '95 - Jan '18)

Source: Sand Hill Econometrics




AI and ML mentions in U.S. Congress (1995-2019)

Source: U.S. Congressional Record website, the McKinsey Global Institute.




Sentiment of articles referencing AI

Source: TrendKite

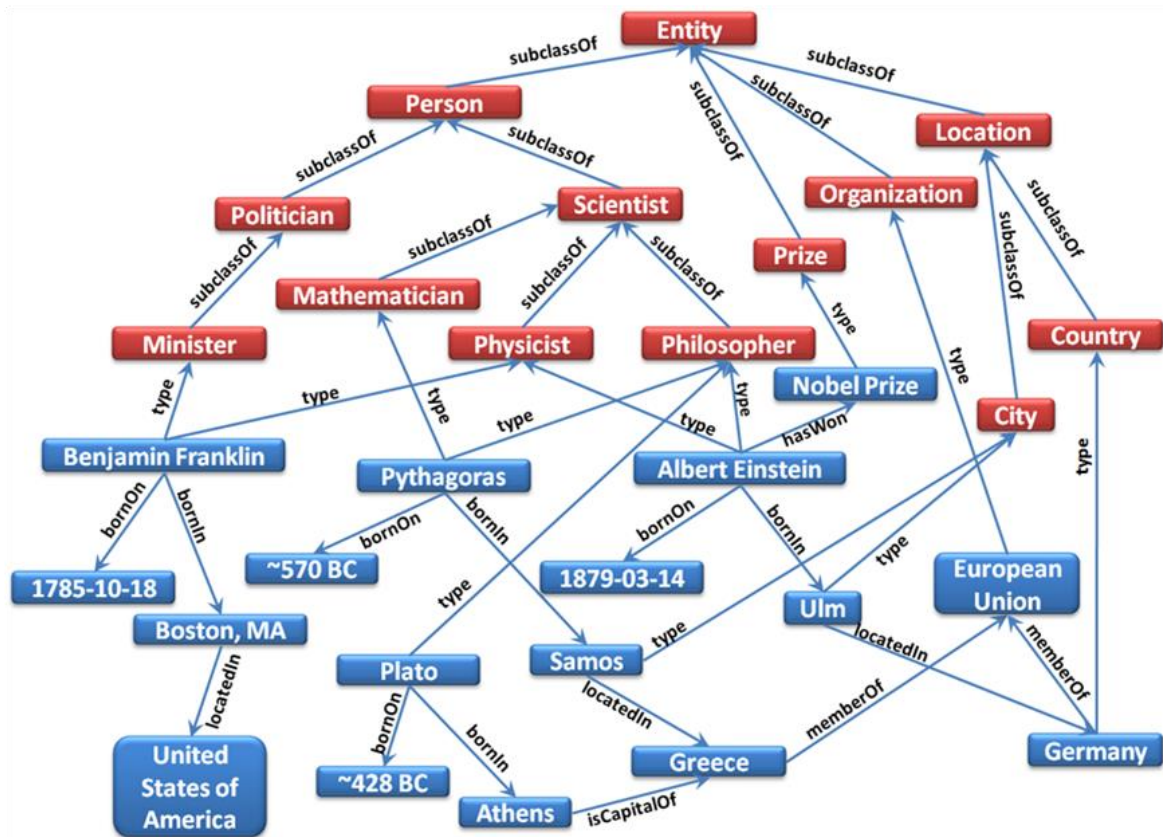


A brief history



Subfields of AI

► Knowledge Representation and Reasoning



Thomas Jefferson

3rd U.S. President

Thomas Jefferson was an American Founding Father, the principal author of the Declaration of Independence, and the third President of the United States. [Wikipedia](#)

Born: April 13, 1743, Shadwell, VA

Died: July 4, 1826, Charlottesville, VA

Presidential term: March 4, 1801 – March 4, 1809

Spouse: [Martha Jefferson](#) (m. 1772–1782)

Party: [Democratic-Republican Party](#)

Awards: AIA Gold Medal

Get updates about Thomas Jefferson

Keep me updated

People also search for

View 15+ more



John Adams



George Washington



Benjamin Franklin



James Madison



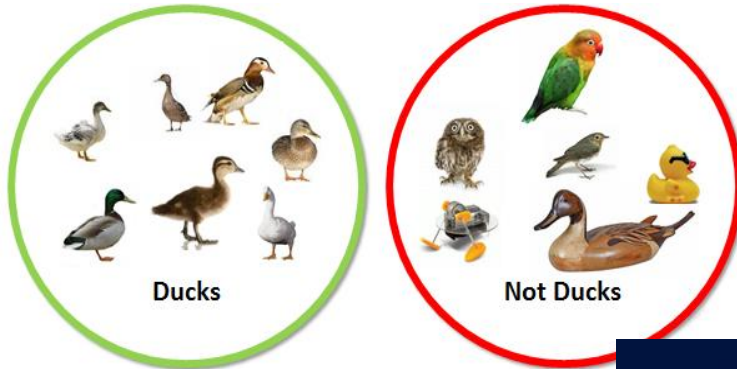


Alexander Hamilton

Feedback

Subfields of AI

► Machine Learning




► Natural Language Processing




Subfields of AI

► Computer Vision



Subfields of AI

► Robotics



Subfields of AI

► Multi-Agent System



Subfields of AI

Integration

Multi-Agent System

Robotics

Natural Language
Processing

Modality-Specific

Computer
Vision

Speech
Recognition

Foundation

Machine
Learning

Knowledge
Representation
& Reasoning

Uncertainty
in AI



Applications of AI

- ▶ Spam email filter
- ▶ Speech recognition
- ▶ Search engine
- ▶ Modern Chinese IME
- ▶ Machine translation



Hi,
Sogou (搜狗) is a popular chinese input software. For Ubuntu 14.04, they worked with the

Google

翻譯

關閉即時翻譯

中文 日文 英文 偵測語言

英文 中文(繁體) 日文

翻譯


The conference will last two days and promises to include the usual mix of executive keynotes, product demos and developer sessions, though we likely won't hear additional details from Facebook until much closer to the event.

本次會議將持續兩天，並承諾執行包括主題演講，產品演示和開發人員的會議通常的搭配，雖然我們可能不會聽到來自Facebook的更多詳細信息，直到更接近事件。

☆ □ Ä 🔊 <


Applications of AI

- ▶ Financial trading
- ▶ Game AI
- ▶ Customer service chatbot
- ▶ Self-driving




Strong AI vs. Weak AI

- ▶ Weak AI (Applied AI)
 - ▶ AI that accomplishes specific tasks



Strong AI vs. Weak AI

- ▶ Strong AI (General AI)
 - ▶ human-like intelligence – AI that could successfully perform any intellectual task that a human can



Strong AI vs. Weak AI




- ▶ Q1: What is the woman in the middle doing?
 - ▶ Action recognition, a CV problem
 - ▶ Q2: What is the woman on the left going to do?
 - ▶ Reasoning about context and intention (beyond current CV)
 - ▶ Q3: This photo was taken in Europe in 2015. What was going on?
 - ▶ Knowing background knowledge by reading news, ...
-



Central problems of (strong) AI


- ▶ Knowledge Representation (KR)
 - ▶ Knowledge: facts, beliefs, concepts, skills, ... that are accumulated over time



Central problems of (strong) AI


- ▶ Inference

- ▶ How to utilize knowledge to derive new information based on existing information



Central problems of (strong) AI


- ▶ Learning
 - ▶ How to accumulate knowledge from experience and education




Three types of approaches

Symbolism

$+$ $-$ \times \div
 \neg \vee \perp \approx
 \in \cap \subseteq Σ
 ∂ ∇ \wedge Π



Connectionism



Statistical Approaches



Symbolism

- ▶ Representing knowledge with symbols and their compositions (expressions)
- ▶ Inference and learning is done by manipulating symbols (e.g., logic)

$$\begin{aligned} \forall x \forall y, Human(x) \wedge Place(y) \wedge At(x, y) \wedge Rain(y) \\ \rightarrow \exists z, Umbrella(z) \wedge Use(x, z) \end{aligned}$$



Symbolism



- ▶ History

- ▶ Dominant during 1950s – 1980s
- ▶ Fell out of favor in 1980s – 1990s
- ▶ Integration with statistical approaches (2000s)
- ▶ Integration with neural approaches (2010s)



Connectionism

- ▶ Representing knowledge with interconnected networks of simple units
 - ▶ Neural networks
- ▶ Inference
 - ▶ Follow the computation specified by the network from input to output
- ▶ Learning
 - ▶ optimization of connection weights




Connectionism

- ▶ History of connectionism: rose and fell for several times
 - ▶ 1940s: pioneer work, e.g., McCulloch-Pitts model
 - ▶ 1958: invention of perceptron (Rosenblatt)
 - ▶ 1969: “Perceptron” published (Minsky & Papert)
 - ▶ Publicized key issues of perceptron (e.g., XOR)
 - ▶ 1970s: AI winter
 - ▶ 1980s: revival of connectionism
 - ▶ Hopfield net, BP algorithm
 - ▶ Rumelhart & McClelland (1986): Parallel Distributed Processing
 - ▶ 1990s-2000s: overtaken in popularity by other methods
 - ▶ 2010s: rise of deep learning
 - ▶ Since ~2012: dominates CV
 - ▶ Since ~2015: dominates NLP




Connectionism



Statistical Approaches

- ▶ Representing knowledge with probabilistic models
- ▶ Inference and learning is done by probabilistic inference




Statistical Approaches

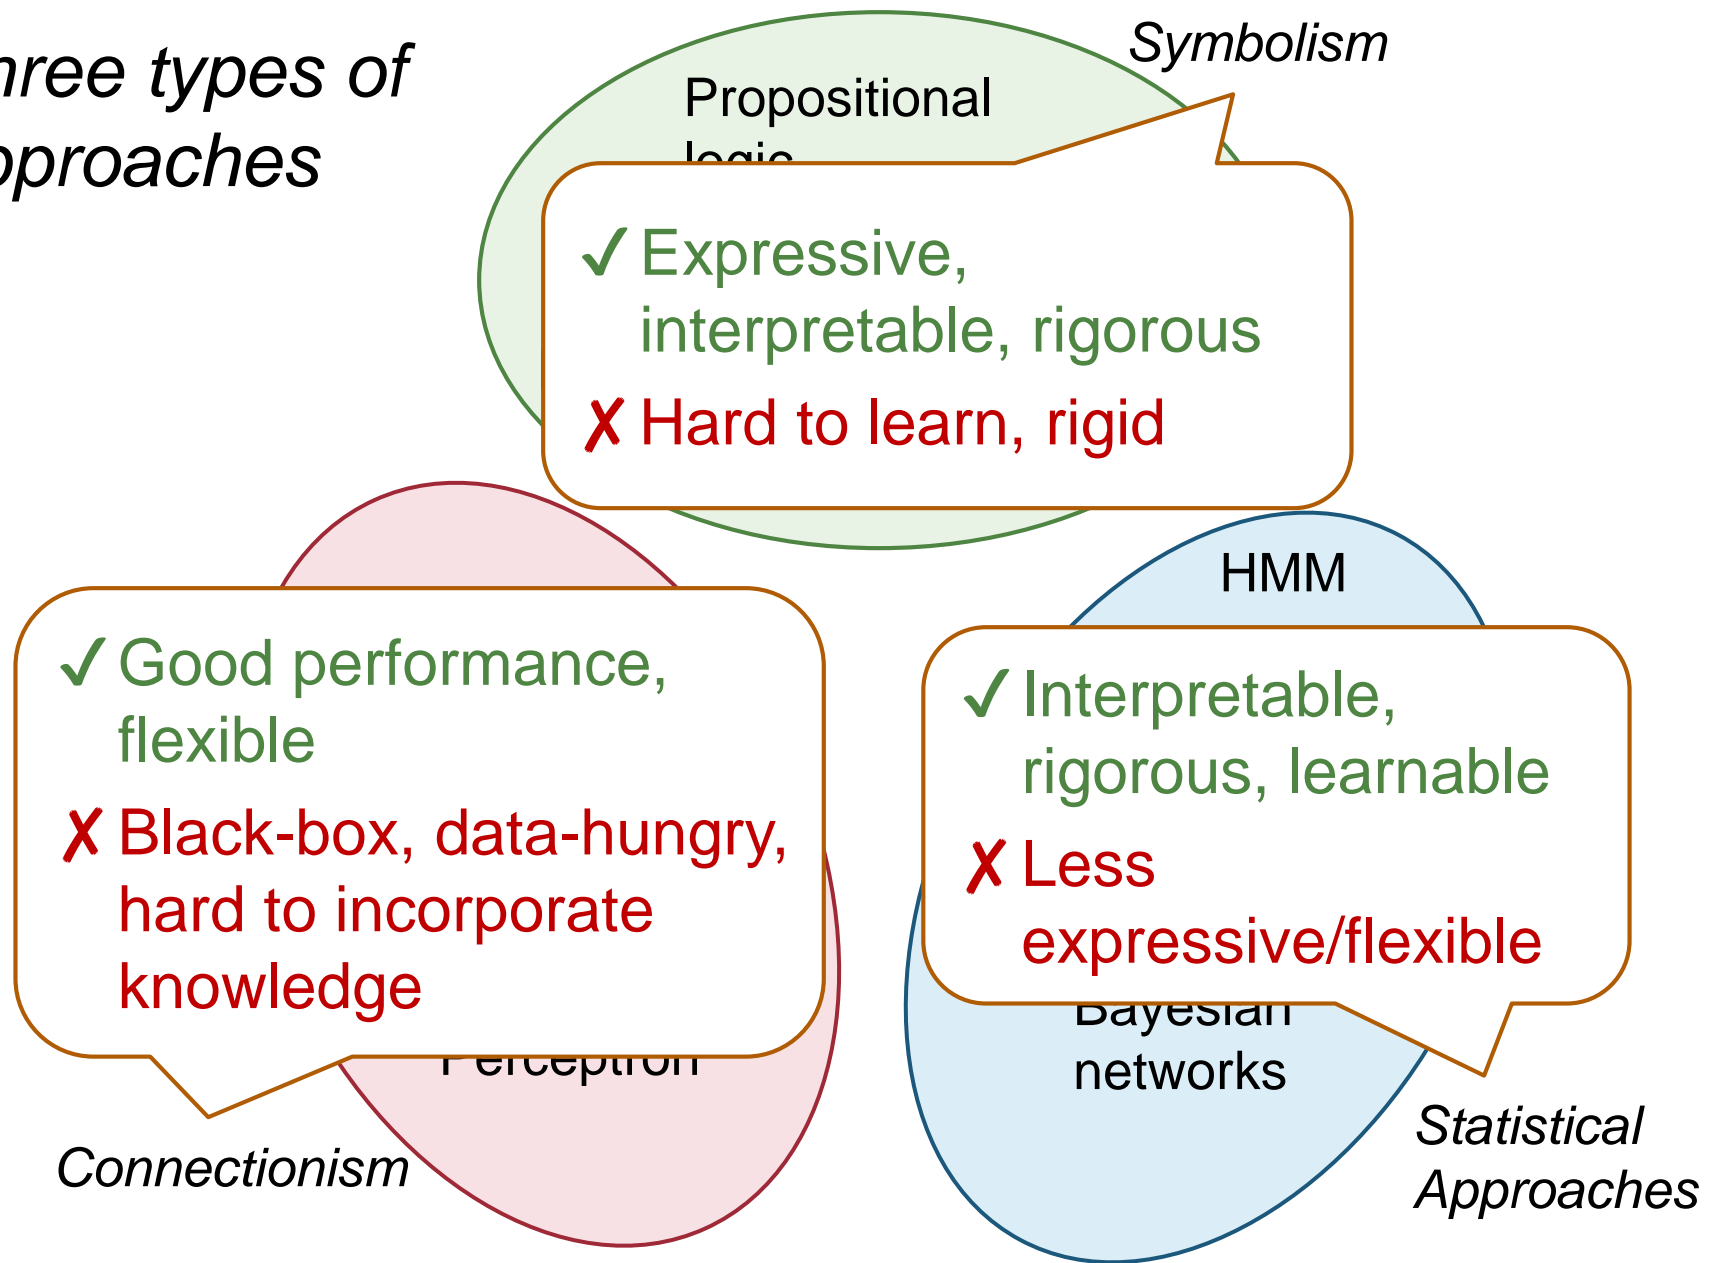
- ▶ History
 - ▶ Become popular since 1990s
 - ▶ Dominant during 2000s
 - ▶ Overshadowed by deep learning in 2010s




Three types of approaches




Three types of approaches



Trends



Trends



Course Overview

Kewei Tu / Xuming He

- ▶ Search
 - ▶ Constraint satisfaction problems
 - ▶ Game
 - ▶ Propositional logic
 - ▶ First-order predicate logic
 - ▶ Probabilistic graphical models
 - ▶ Probabilistic temporal models
 - ▶ Probabilistic logics
 - ▶ Introduction to natural language processing
 - ▶ Introduction to machine learning
 - ▶ Markov decision processes
 - ▶ Reinforcement learning
-

