


Announcement

- ▶ Homework 3
 - ▶ Due: Nov. 11, 11:59pm
- ▶ Programming Assignment 3
 - ▶ Due: Nov. 25, 11:59pm (three weeks)



Three types of (strong) AI approaches



Probabilistic Logics

AIMA 14.6
Additional materials

Additional reference materials

- ▶ L. Getoor and B. Taskar (eds.), Introduction to Statistical Relational Learning, 2007. Cambridge, MA: MIT Press.
 - ▶ Ch 5: Probabilistic Relational Models
 - ▶ Ch 12: Markov Logic



Logics vs. Probabilistic Models

- ▶ Symbolic logics
 - ▶ FOL is very expressive
 - ▶ relations between objects, quantifiers
 - ▶ But it cannot model uncertainty
- ▶ Probabilistic Models
 - ▶ BN/MN model uncertainty in a concise manner
 - ▶ But limited in expressiveness
 - ▶ BN/MN is essentially propositional



Probabilistic Logics

- ▶ Goal
 - ▶ Combine (subsets of) logic and probability into a single language
- ▶ A.k.a. Statistical Relational Learning
- ▶ Lots of approaches. We will cover two of them:
 - ▶ Probabilistic Relational Models
 - ▶ Markov Logic





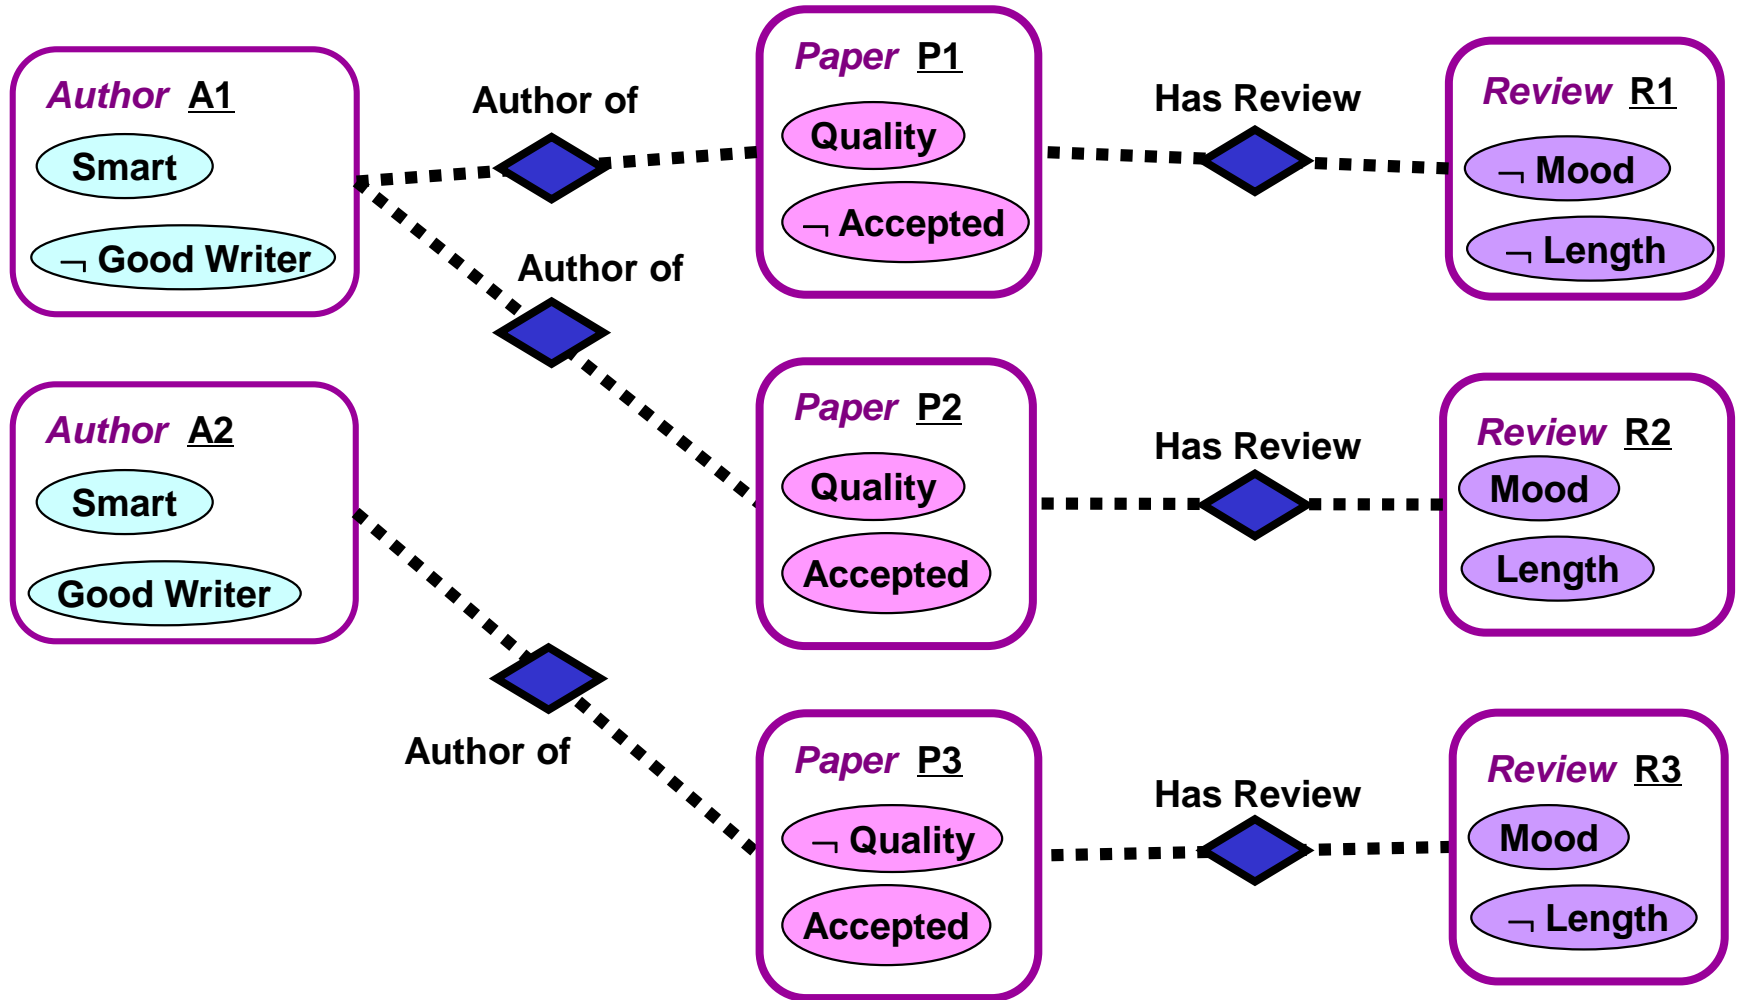
Probabilistic Relational Models

Probabilistic Relational Models

- ▶ Logical language
 - ▶ Frame (typed relational knowledge)
 - ▶ A subclass of FOL
- ▶ Probabilistic language
 - ▶ Bayes nets

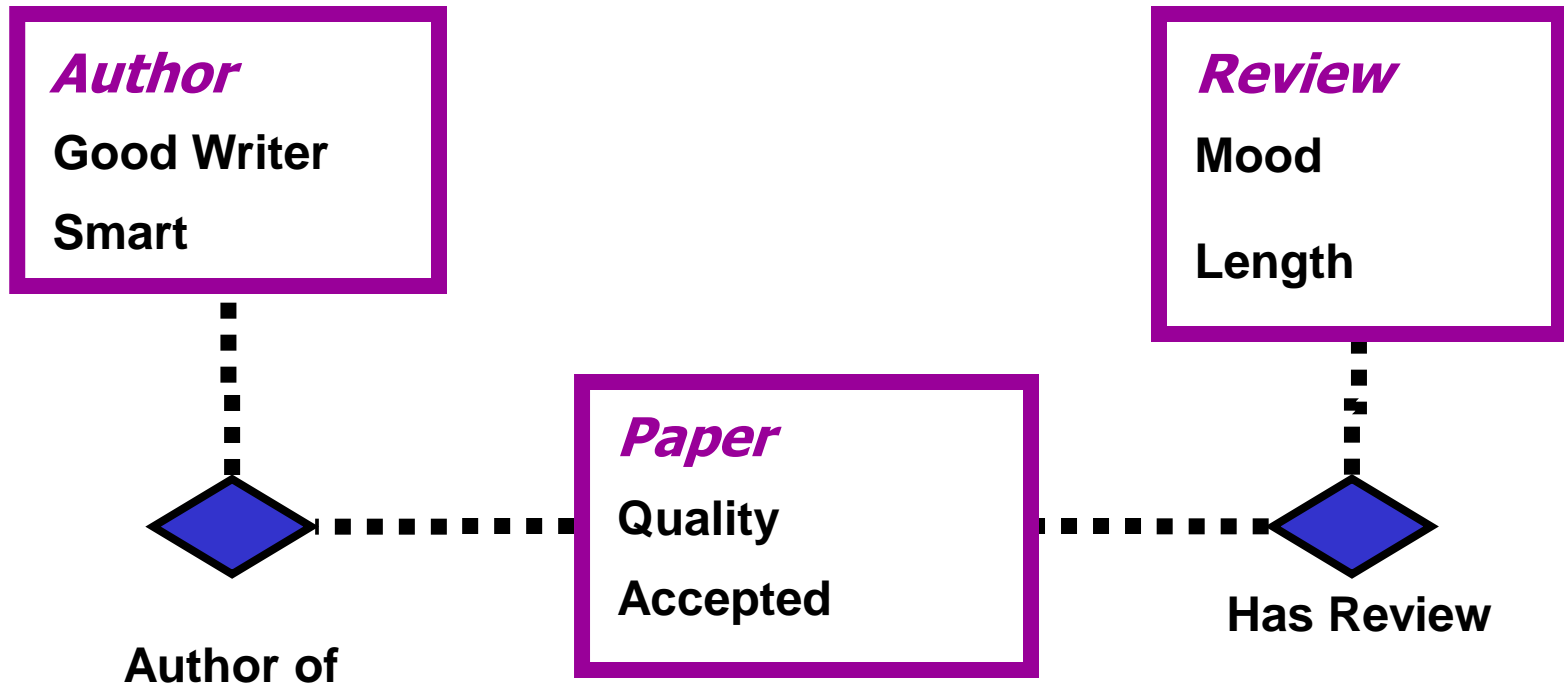


Typed relational knowledge



Typed relational knowledge

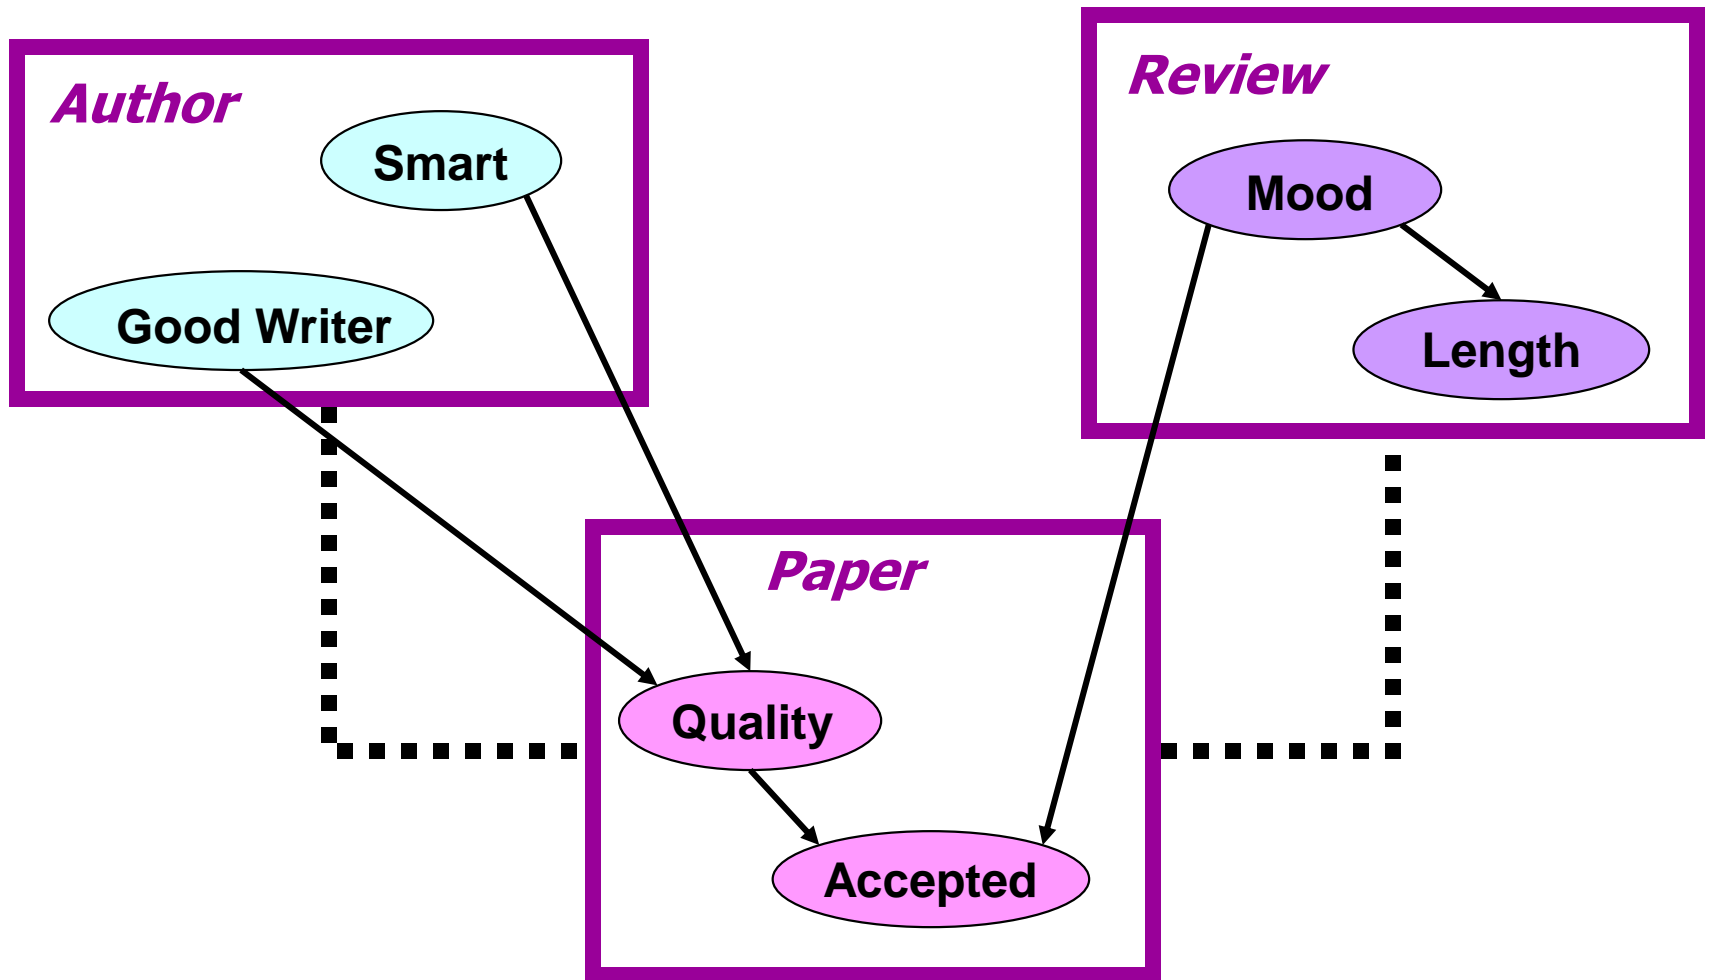
Why is this a subclass of FOL?



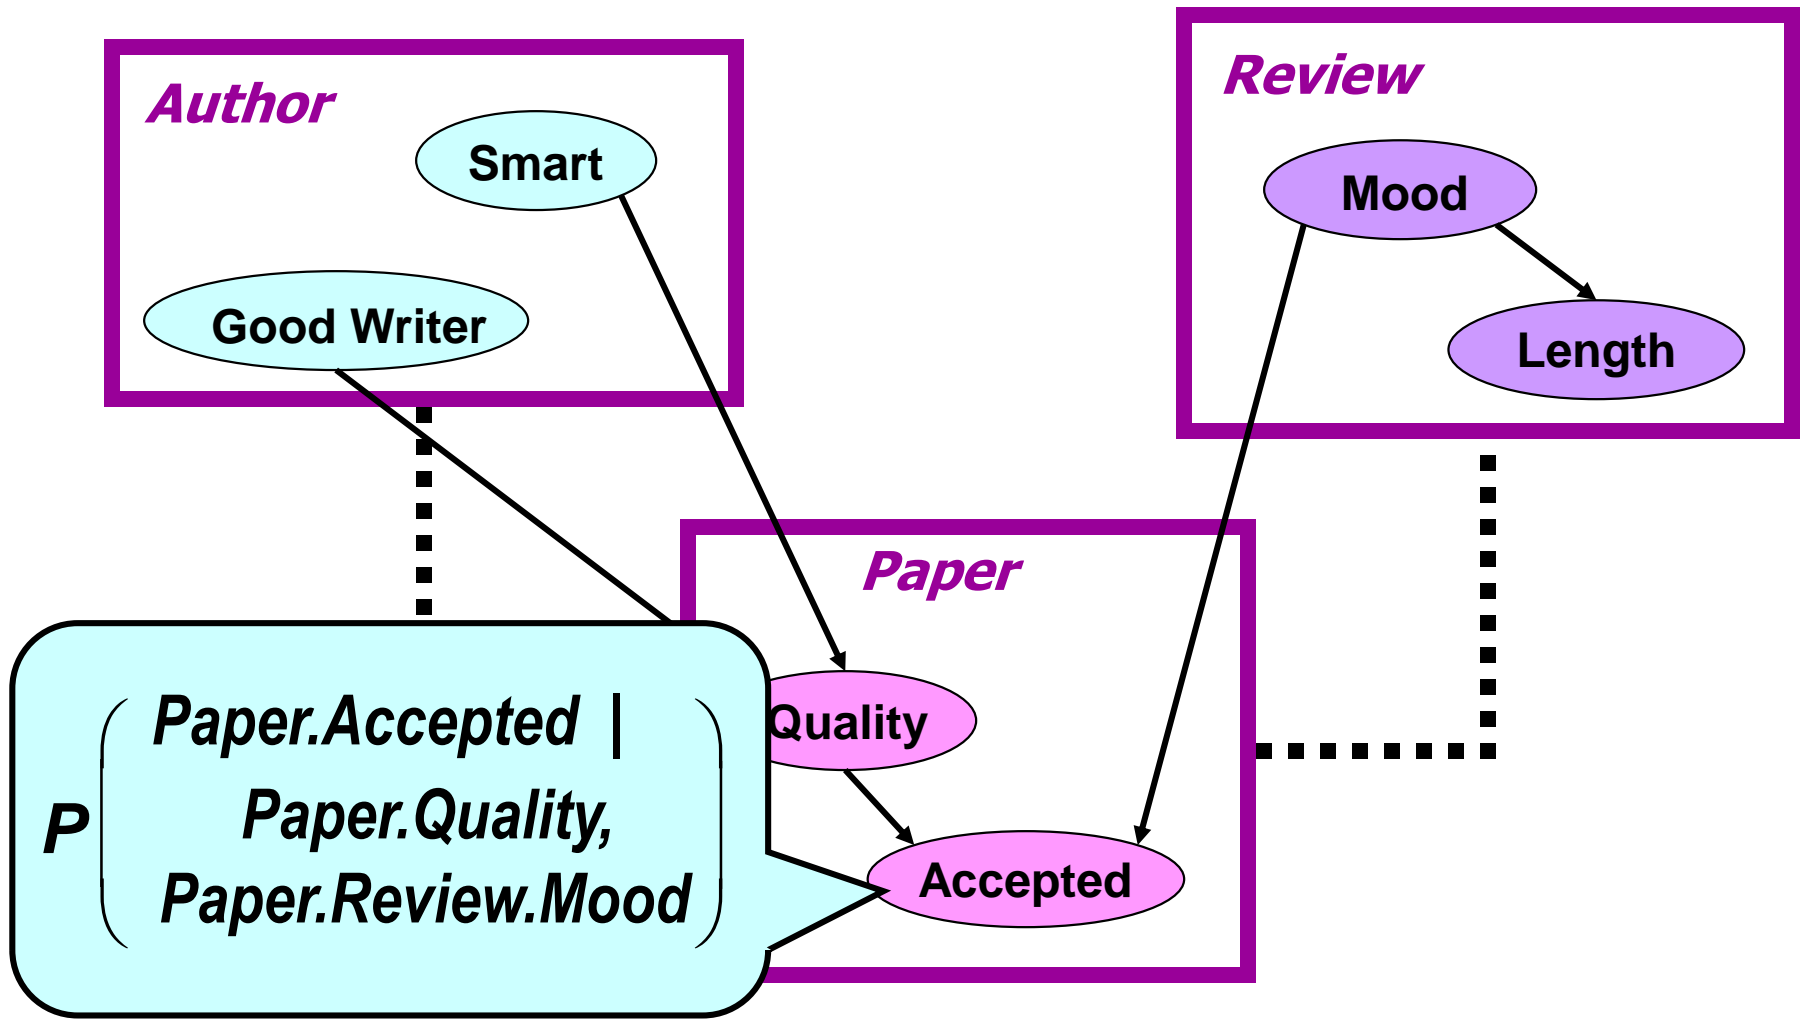
Ontology / Schema

- ▶ The types of objects and their valid relations and attributes

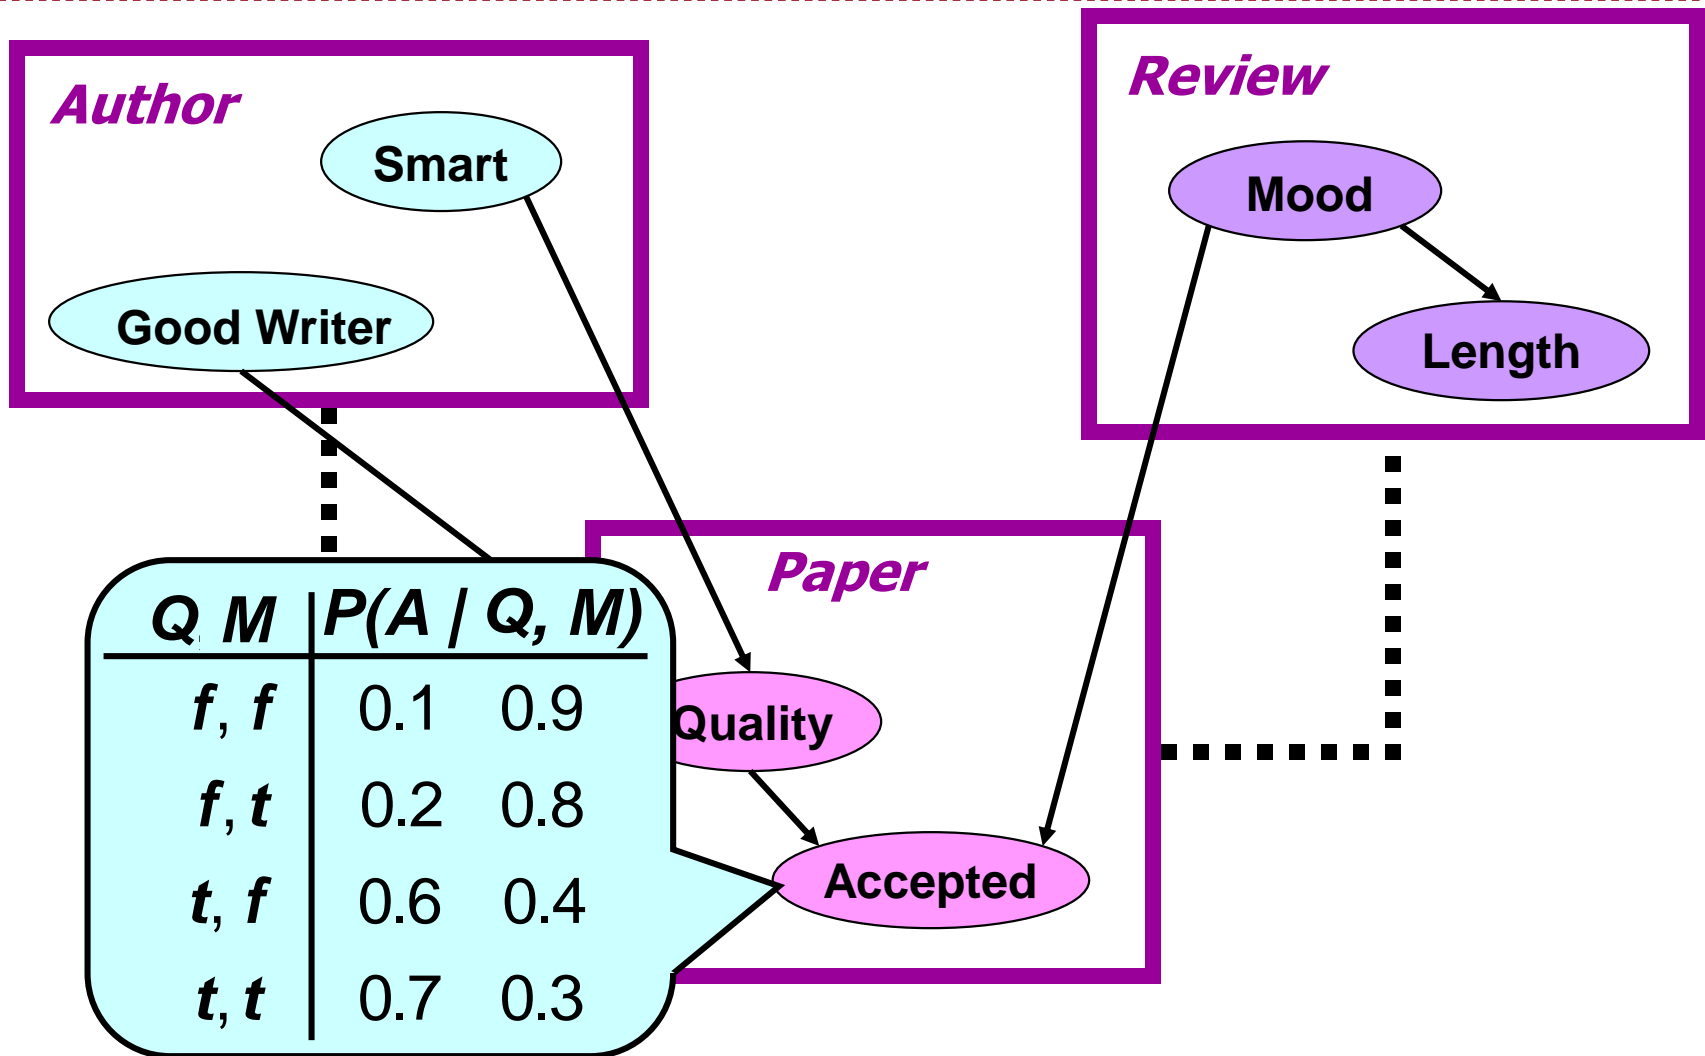
Probabilistic Relational Model




Probabilistic Relational Model



Probabilistic Relational Model



Relational Skeleton

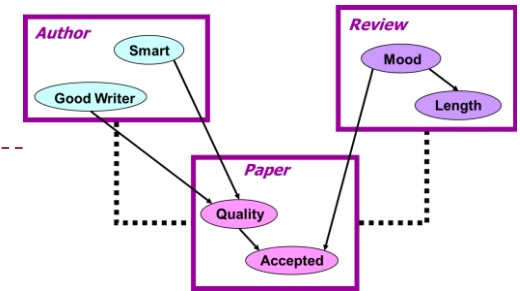
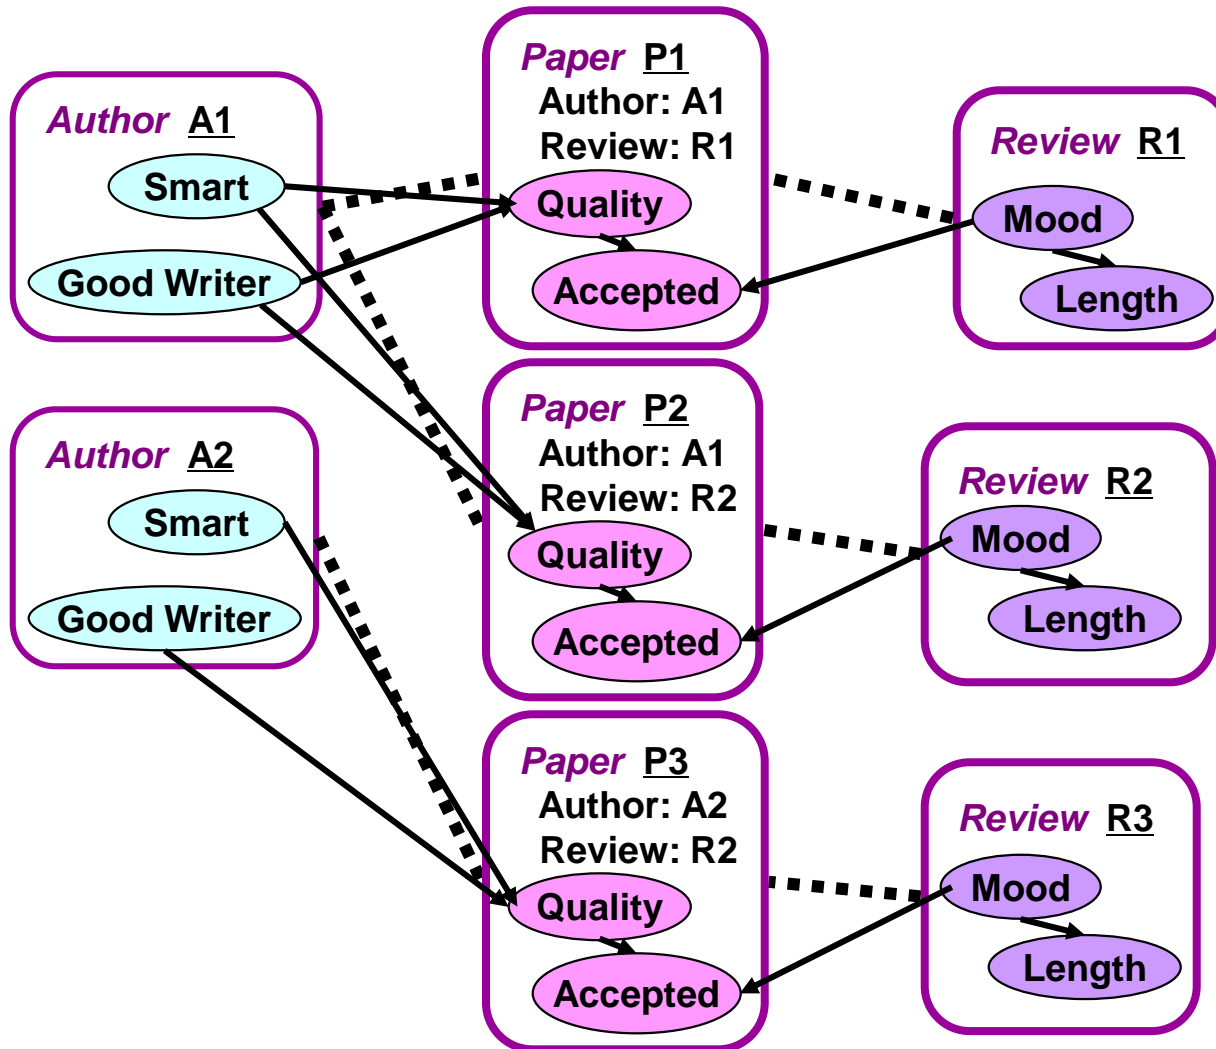


Fixed relational skeleton σ :

- set of objects in each class
- relations between them
- attribute values unknown




PRM with Attribute Uncertainty

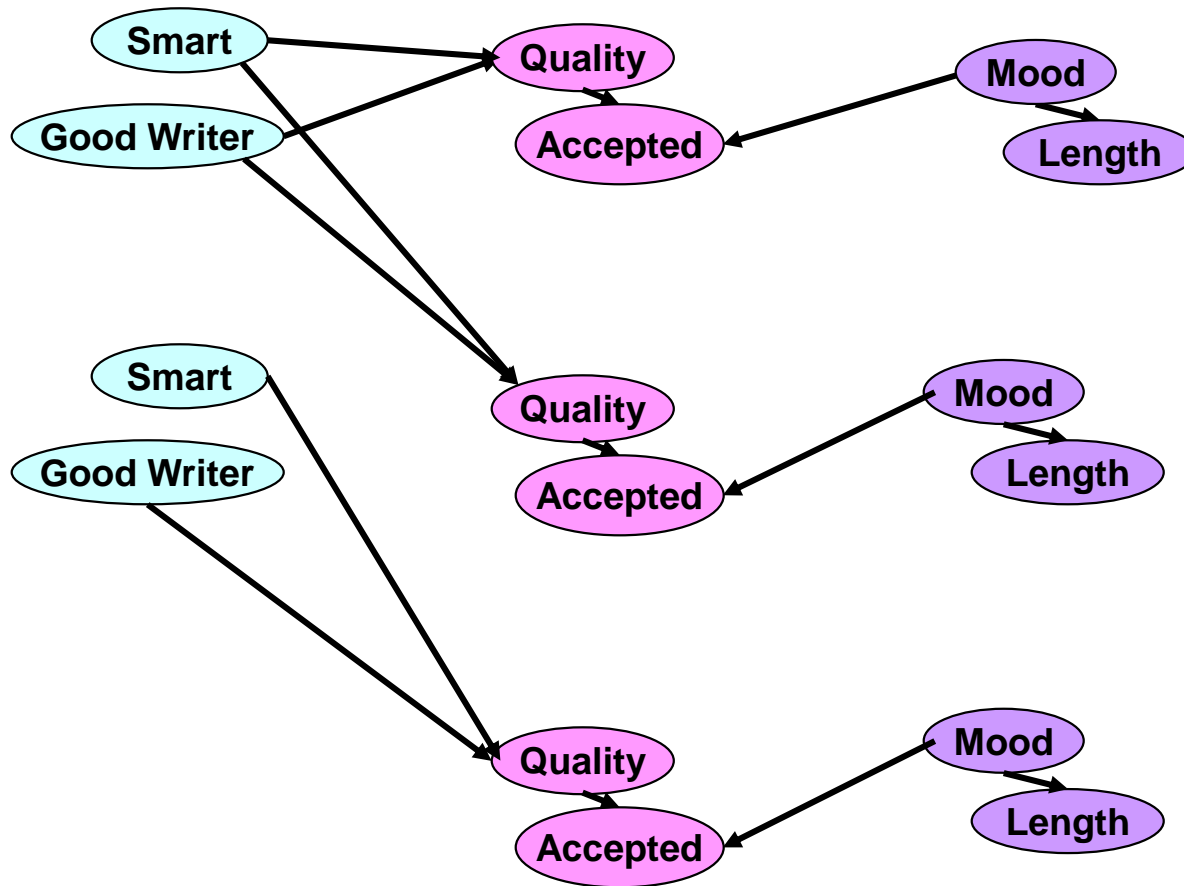


PRM **unrolled** wrt. the relational skeleton produces a BN that models the distribution over instantiations of attributes.

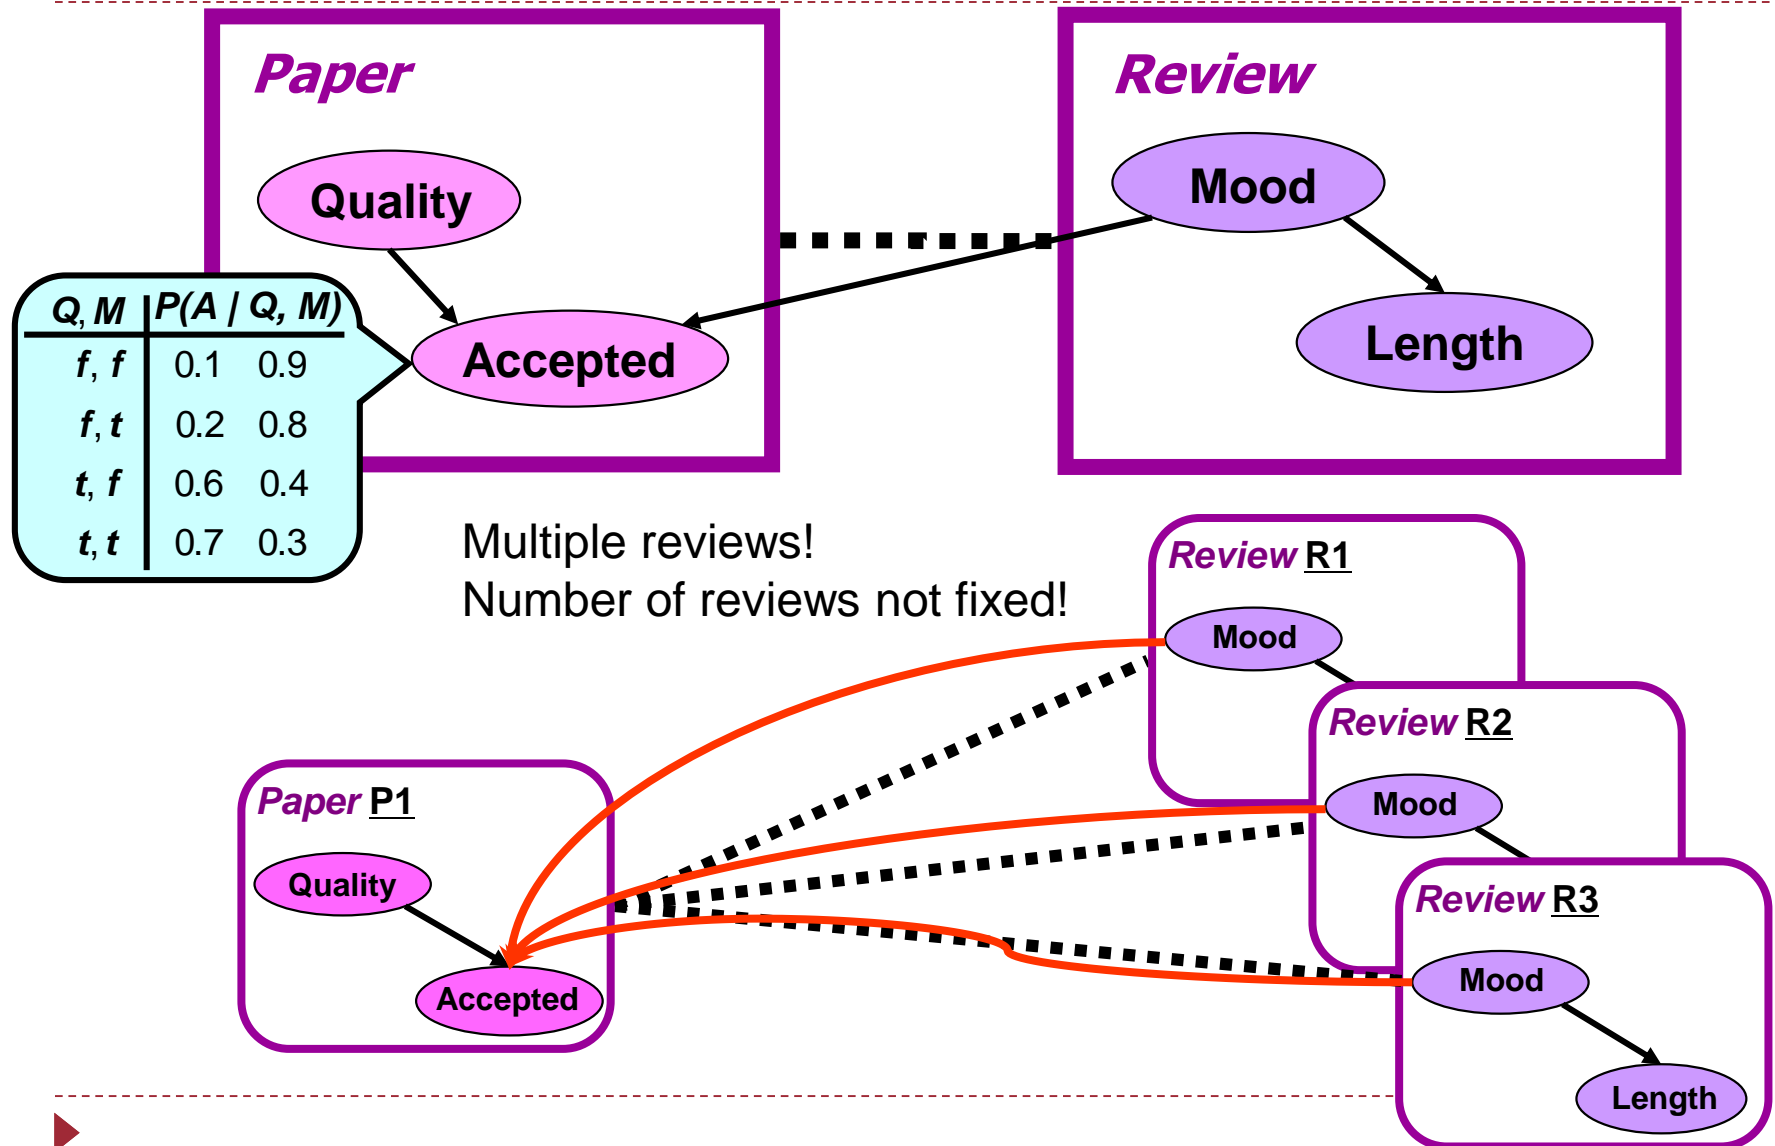
PRM with Attribute Uncertainty



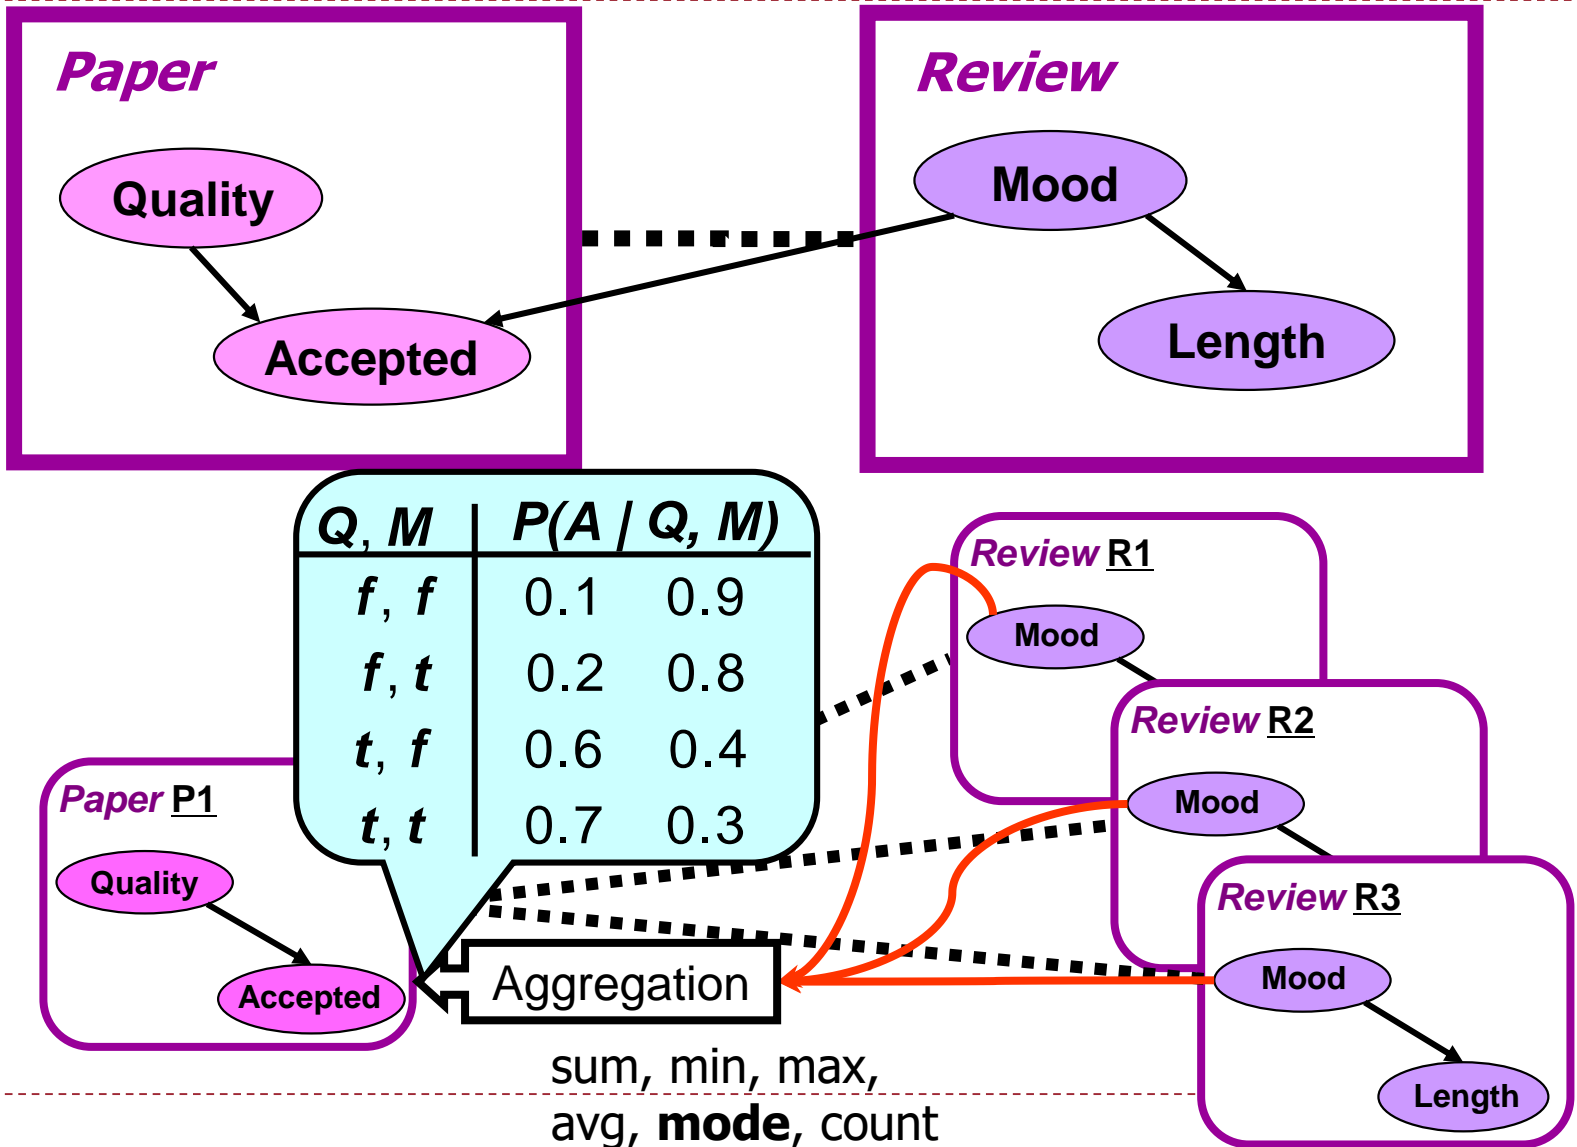
PRM **unrolled** wrt. the relational skeleton produces a BN that models the distribution over instantiations of attributes.





PRM: Aggregate Dependencies



PRM: Aggregate Dependencies



PRM with Attribute Uncertainty



PRM (\mathcal{S}, Θ) + relational skeleton (σ) =

probability distribution over instantiations of attributes \mathcal{I} :

$$P(\mathcal{I} \mid \sigma, \mathcal{S}, \Theta) = \prod_{x \in \sigma} \prod_{x.A} P(x.A \mid \text{parents}_{\mathcal{S}, \sigma}(x.A))$$

\nearrow
Objects

\nwarrow
Attributes



Structural Uncertainty

- ▶ PRM with AU only well-defined when the relational skeleton is known
- ▶ What if we are uncertain about the relational structure?
 - ▶ How many objects does an object relate to?
 - ▶ Which object is an object related to?
 - ▶ Does an object actually exist?
 - ▶ Are two objects identical?




Structural Uncertainty


- ▶ Need probabilistic models that capture **structural uncertainty**
- ▶ Types of SU:
 - ▶ Existence uncertainty
 - ▶ Reference uncertainty
 - ▶ Number uncertainty
 - ▶ Type uncertainty
 - ▶ Identity uncertainty



Citation Schema



Existence Uncertainty




PRM with Existence Uncertainty



Introduce the **Exists** attribute for *Cites*




PRM with Existence Uncertainty



Citer.Topic	Cited.Topic	False	True
<i>Theory</i>	<i>Theory</i>	0.995	0.005
<i>Theory</i>	<i>AI</i>	0.999	0.001
<i>AI</i>	<i>Theory</i>	0.997	0.003
<i>AI</i>	<i>AI</i>	0.992	0.008

Object skeleton

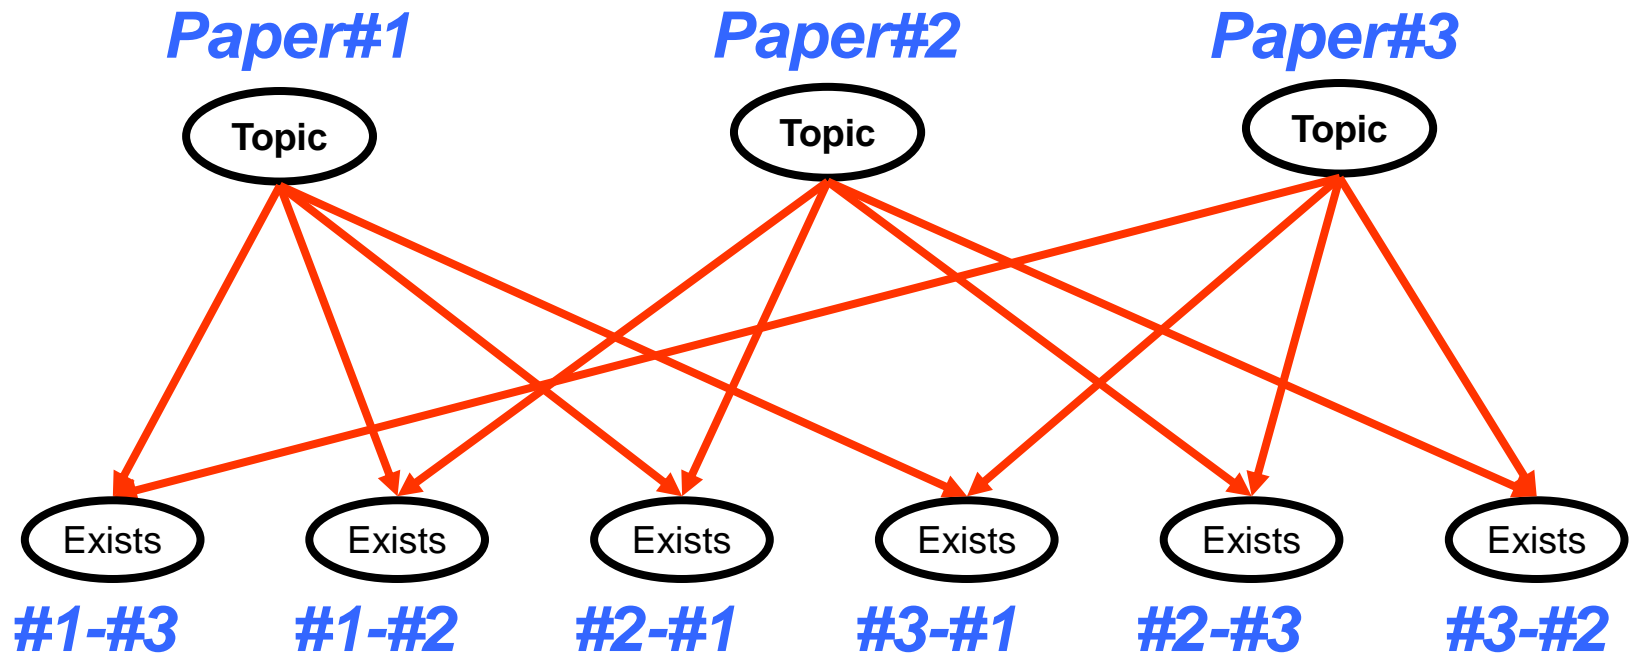


Fixed object skeleton σ :

- set of objects in each class
- unknown relations between them
- unknown attribute values




PRM with Existence Uncertainty

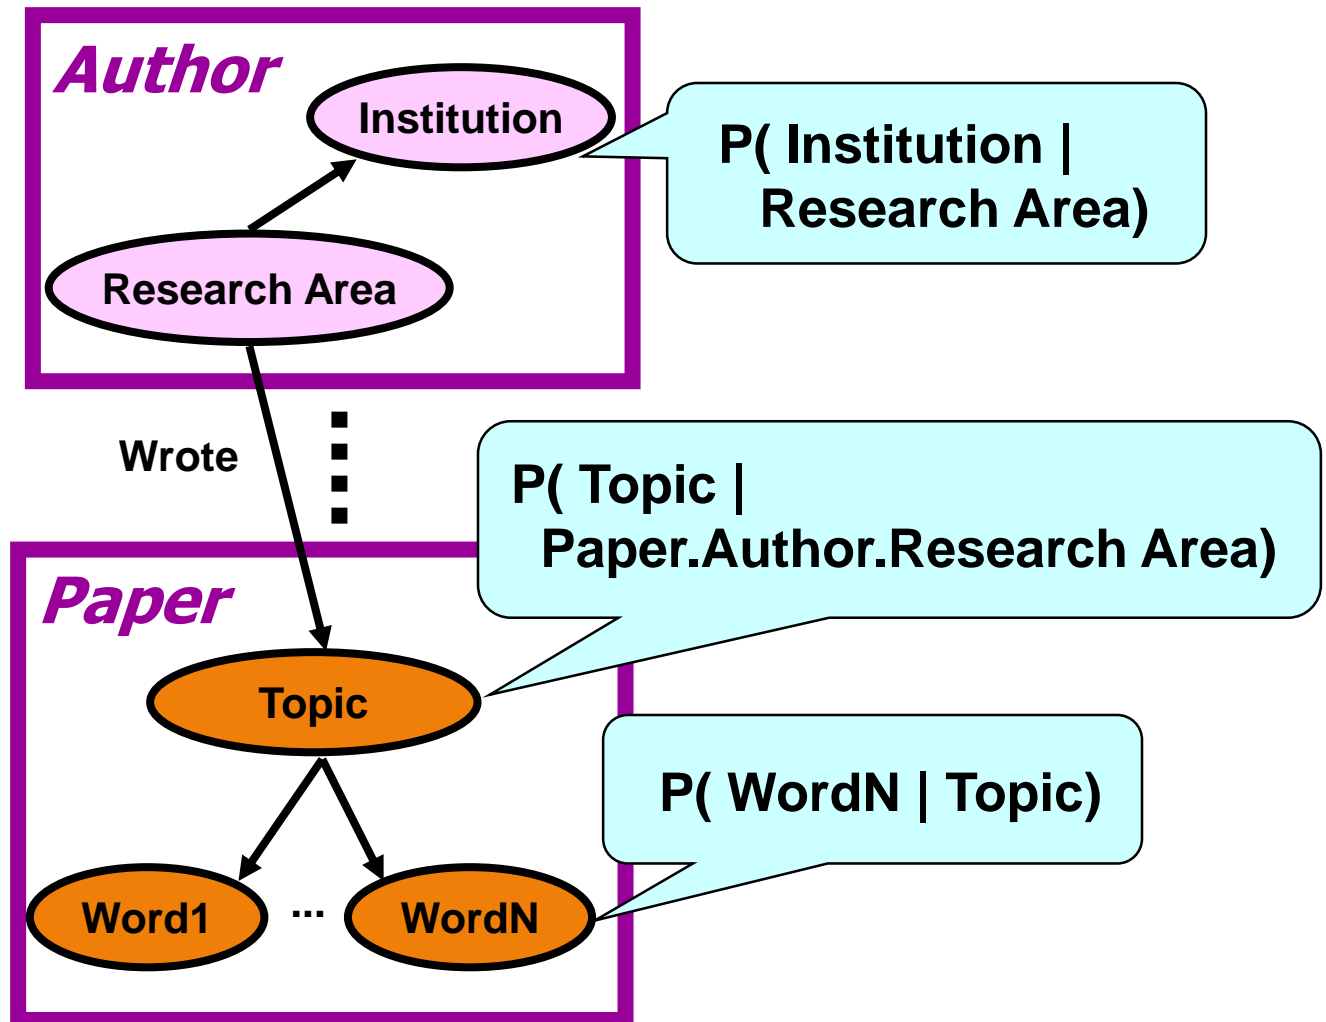


PRM w/ EU unrolled wrt. the object skeleton
produces a BN

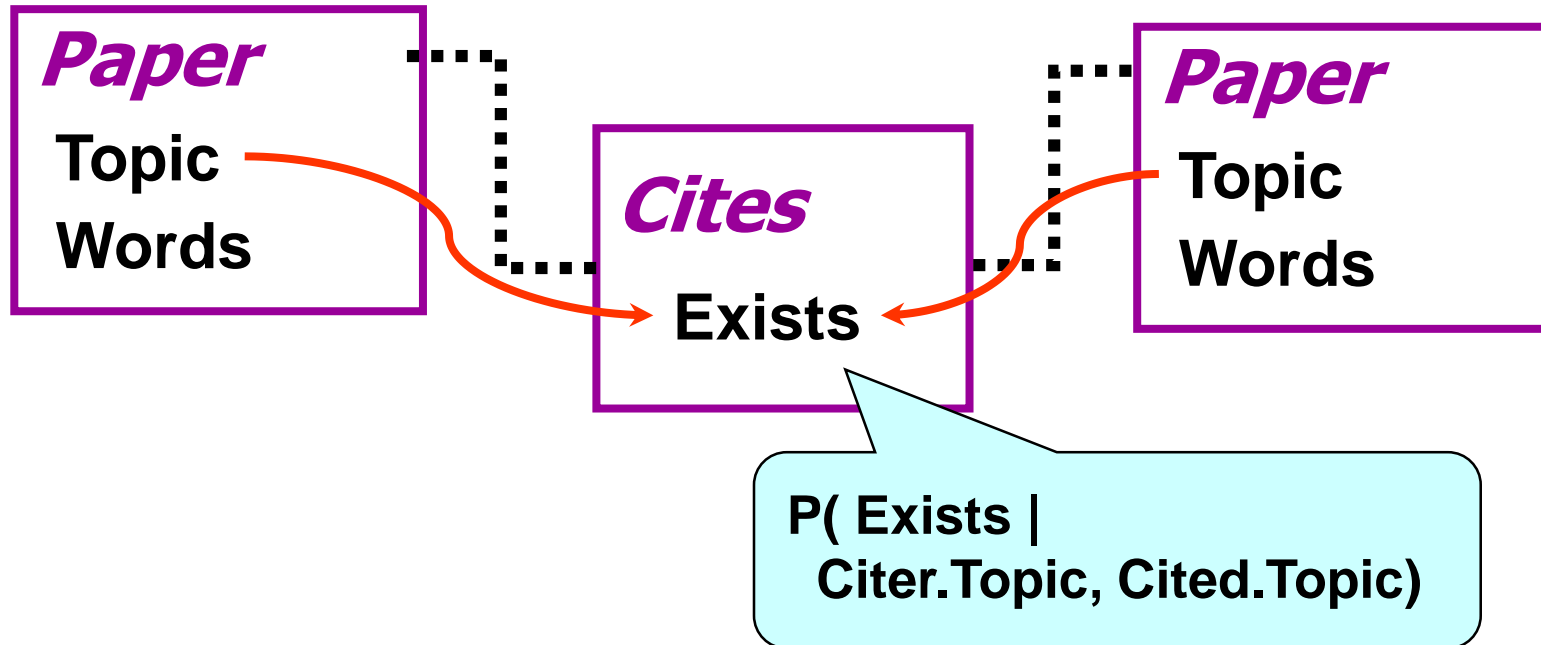
A more complicated example



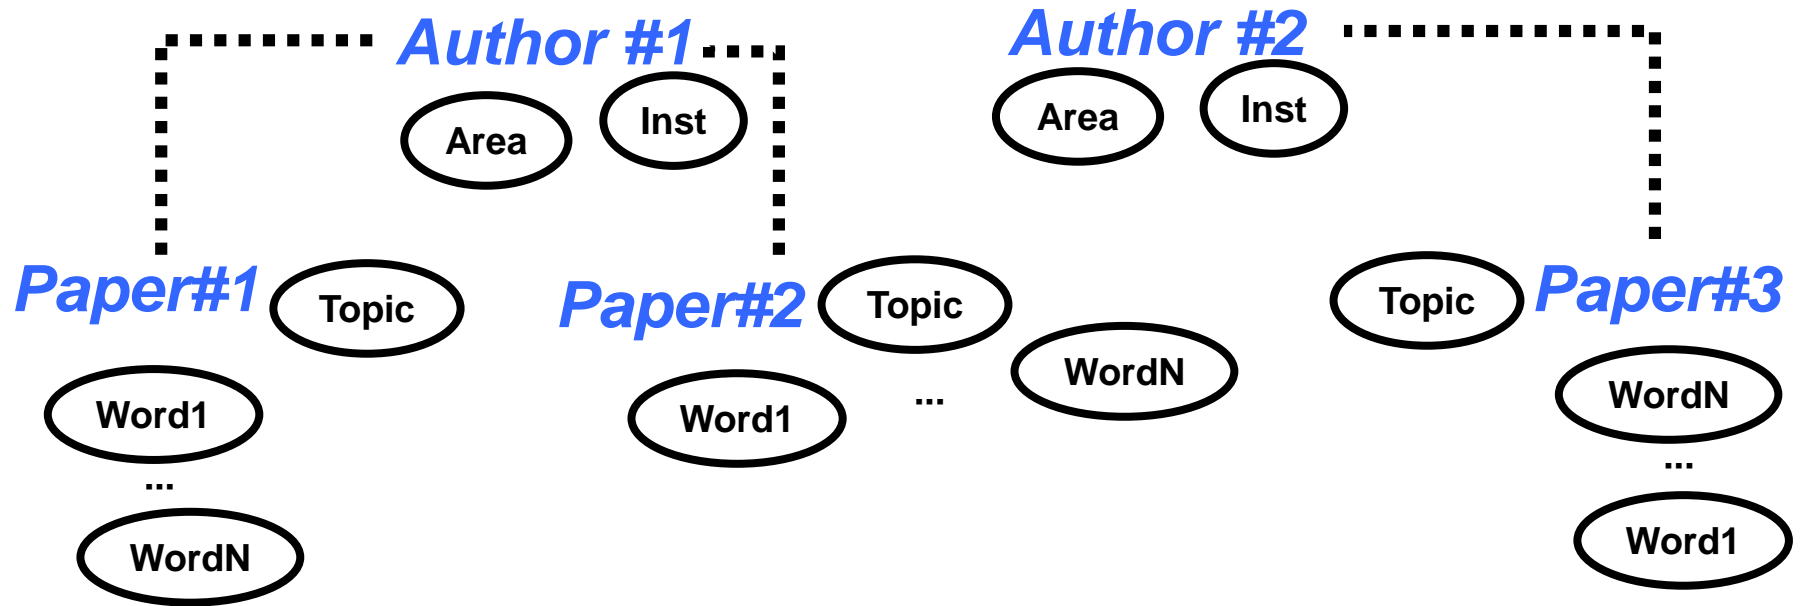
PRM with Attribute Uncertainty



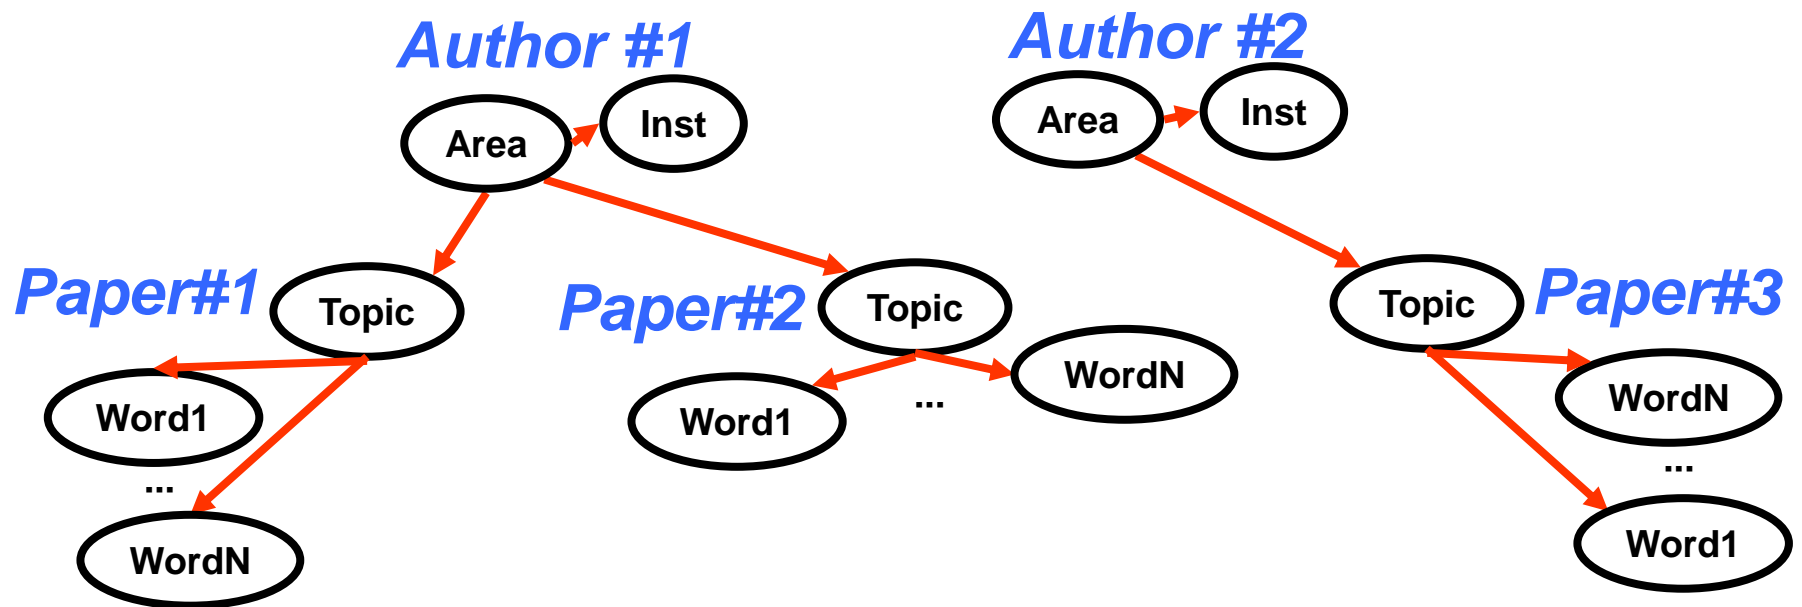
PRM with Existence Uncertainty



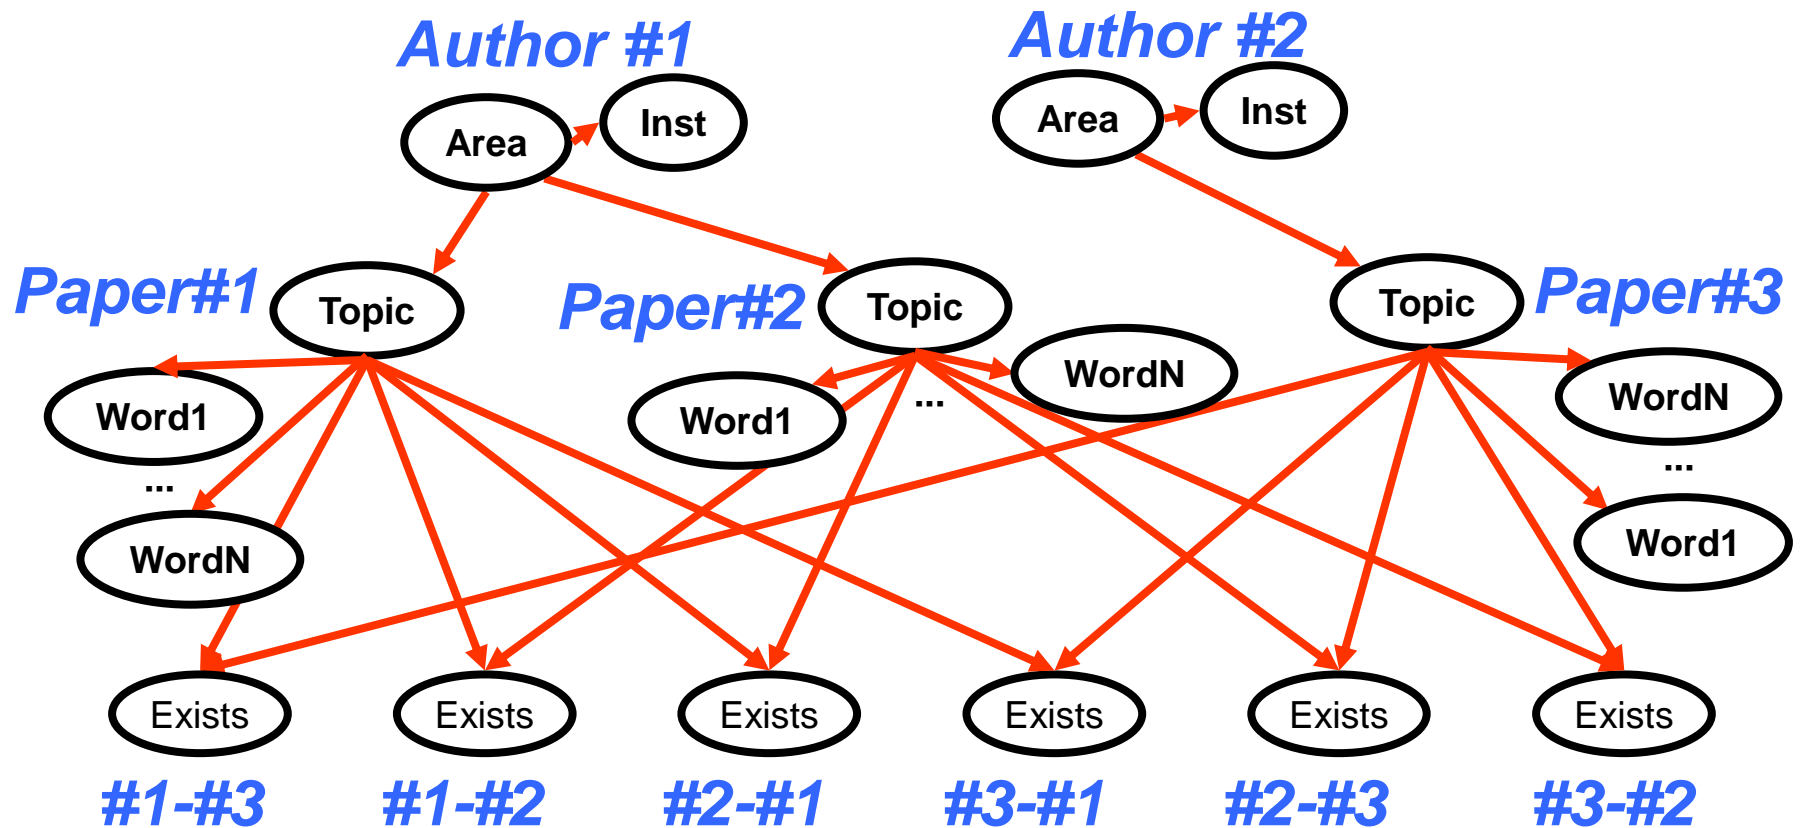
Unrolling



Unrolling



Unrolling





Markov Logic

Markov Logic

- ▶ Logical language
 - ▶ First-order logic
- ▶ Probabilistic language
 - ▶ Markov networks



Review: Markov networks


- ▶ A Markov network (or Markov random field) encodes a joint distribution with an undirected graph

$$p(\mathbf{x}) = \frac{1}{Z} \prod_C \psi_C(\mathbf{x}_C)$$

where $\psi_C(\mathbf{x}_C)$ is the **potential** over **clique** C and

$$Z = \sum_{\mathbf{x}} \prod_C \psi_C(\mathbf{x}_C)$$

is the **normalization coefficient**.



Markov Logic: Intuition

- ▶ A logical KB is a set of **hard constraints** on the set of possible worlds
 - ▶ If a world violates a formula, it becomes impossible
- ▶ Let's make them **soft constraints**: When a world violates a formula, it becomes less probable, not impossible
- ▶ Give each formula a **weight** (Higher weight \Rightarrow Stronger constraint)

$$P(\text{world}) \propto \exp\left(\sum \text{weights of formulas it satisfies}\right)$$



Markov Logic: Definition

- ▶ A Markov Logic Network (MLN) is a set of pairs (F, w) where
 - ▶ F is a formula in first-order logic
 - ▶ w is a real number



Example: Friends & Smokers

Smoking causes cancer.

Friends have similar smoking habits.



Example: Friends & Smokers

1.5	$\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$
1.1	$\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftrightarrow \text{Smokes}(y))$



Markov Logic: Definition

- ▶ A Markov Logic Network (MLN) is a set of pairs (F, w) where
 - ▶ F is a formula in first-order logic
 - ▶ w is a real number
- ▶ Together with a set of constants, it defines a Markov network with
 - ▶ One node for each grounding of each predicate in the MLN
 - ▶ This is *propositionalization* (remember?)



Example: Friends & Smokers

1.5 $\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$

1.1 $\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftrightarrow \text{Smokes}(y))$

Two constants: **Anna** (A) and **Bob** (B)

Friends(A,B)

Friends(A,A)

Smokes(A)

Smokes(B)

Friends(B,B)

Cancer(A)

Friends(B,A)

Cancer(B)



Markov Logic: Definition

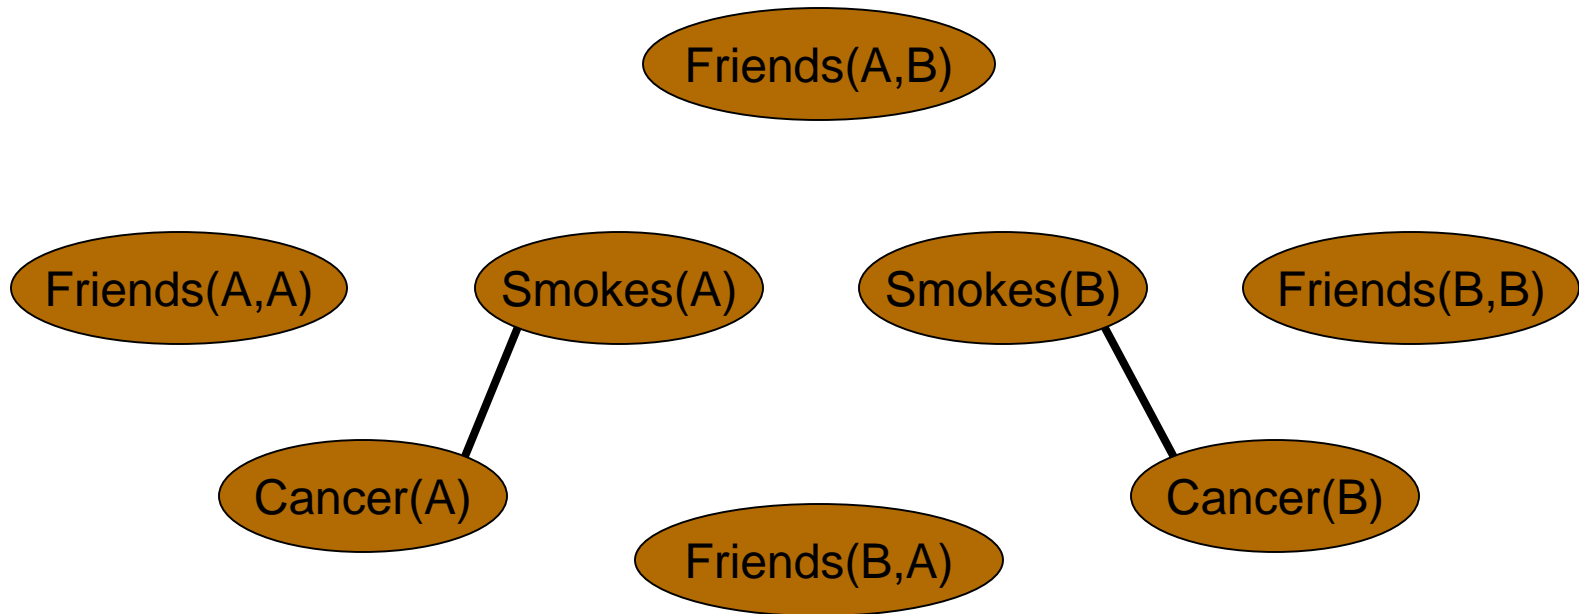
- ▶ A Markov Logic Network (MLN) is a set of pairs (F, w) where
 - ▶ F is a formula in first-order logic
 - ▶ w is a real number
- ▶ Together with a set of constants, it defines a Markov network with
 - ▶ One node for each grounding of each predicate in the MLN
 - ▶ This is *propositionalization* (remember?)
 - ▶ One clique for each grounding of each formula F in the MLN, with the potential being:
 - ▶ $\exp(w)$ for node assignments that satisfy F
 - ▶ 1 otherwise



Example: Friends & Smokers

1.5	$\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$
1.1	$\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftrightarrow \text{Smokes}(y))$

Two constants: **Anna** (A) and **Bob** (B)




Example: Friends & Smokers

1.5 $\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$

1.1 $\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftrightarrow \text{Smokes}(y))$

Two constants: **Anna** (A) and **Bob** (B)




Example: Friends & Smokers

1.5 $\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$

1.1 $\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftarrow$

S	C	$\Phi(S,C)$
False	False	exp(1.5)
False	True	exp(1.5)
True	False	1
True	True	exp(1.5)

Two constants: **Anna** (A) and **Bob** (B)



Markov Logic Networks

- ▶ MLN is **template** for ground Markov nets
- ▶ Probability of a world x :

$$P(x) = \frac{1}{Z} \exp \left(\sum_i w_i n_i(x) \right)$$

Weight of formula i

No. of true groundings of formula i in x

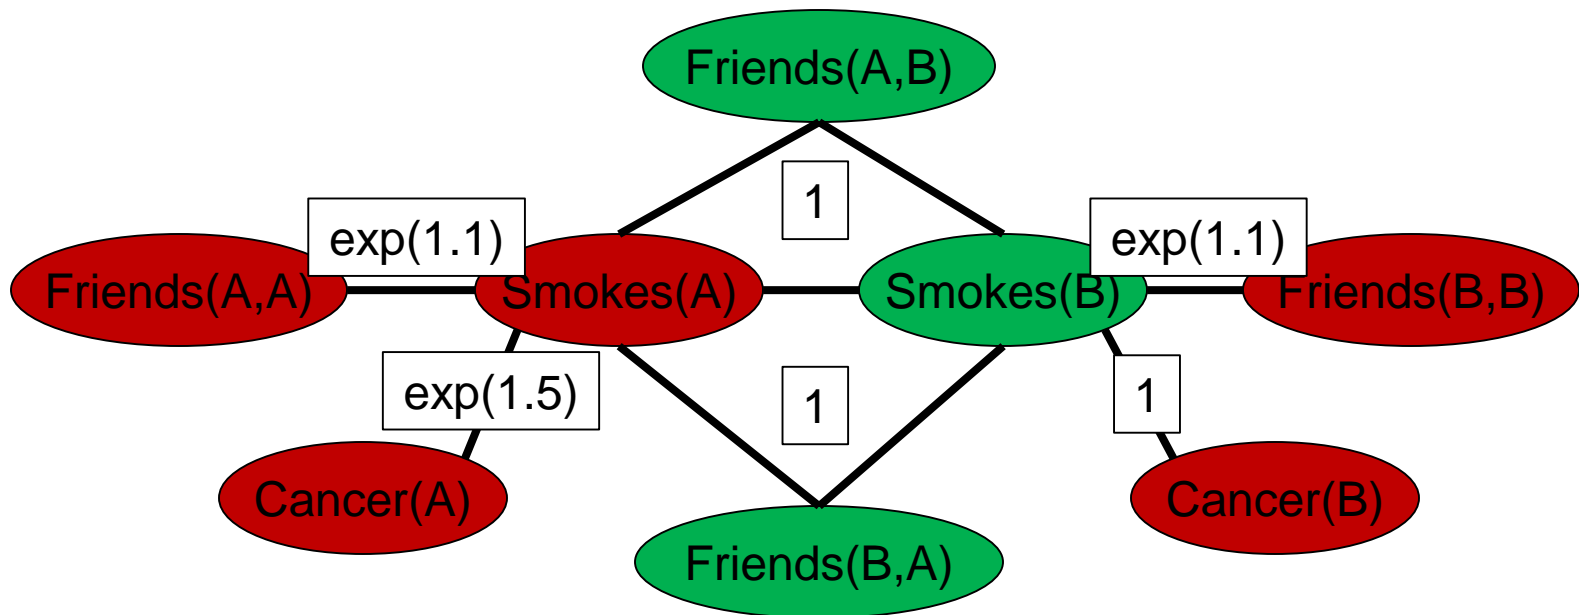


Example: Friends & Smokers

1.5 $\forall x \text{ Smokes}(x) \Rightarrow \text{Cancer}(x)$

1.1 $\forall x, y \text{ Friends}(x, y) \Rightarrow (\text{Smokes}(x) \Leftrightarrow \text{Smokes}(y))$

Two constants: **Anna** (A) and **Bob** (B)



$$P(x) \propto \exp(1.1 + 1.1 + 1.5)$$


Relation to First-Order Logic

- ▶ Infinite weights \Rightarrow First-order logic
 - ▶ $P(x) > 0$ iff. x satisfies KB
- ▶ Markov logic allows contradictions between formulas



Relation to PRM

- ▶ MLN is More general and flexible than PRM
- ▶ In principle, a PRM can be converted into a MLN by writing a formula for each entry of each CPT and setting the weight to be the logarithm of the conditional probability



Q, M	$P(A \mid Q, M)$	
f, f	0.1	0.9
f, t	0.2	0.8
t, f	0.6	0.4
t, t	0.7	0.3
	t	f



$$\log 0.1: \forall x, y \text{ hasReview}(x, y) \wedge Q(x, f) \wedge M(y, f) \wedge A(x, t)$$

$$\log 0.9: \forall x, y \text{ hasReview}(x, y) \wedge Q(x, f) \wedge M(y, f) \wedge A(x, f)$$

$$\log 0.2: \forall x, y \text{ hasReview}(x, y) \wedge Q(x, f) \wedge M(y, t) \wedge A(x, t)$$

.....



Inference

Inference

- ▶ A naive approach
 - ▶ Unroll the model to a BN or MN and run inference algorithms (such as VE)
 - ▶ Problem: the BN/MN may be very large and highly interconnected
- ▶ Lifted inference
 - ▶ Lots of repeated structures in the unrolled model \Rightarrow repeated computation in inference
 - ▶ Group similar random variables at the FOL level and handle them at the same time



Summary

- ▶ Probabilistic Relational Models
 - ▶ Logical language: Frame
 - ▶ Probabilistic language: Bayes nets
 - ▶ Bayes net template for object classes
 - ▶ Object's attrs. can depend on attrs. of related objs.
- ▶ Markov Logic
 - ▶ Logical language: First-order logic
 - ▶ Probabilistic language: Markov networks
 - ▶ Syntax: First-order formulas with weights
 - ▶ Semantics: Templates for Markov net cliques

