

# Natural Language Processing

AIMA Ch 23

# Additional Reference

---

- ▶ [SLP] Speech and Language Processing, Daniel Jurafsky and James H. Martin
  - ▶ 2nd edition, 2008
  - ▶ 3rd edition, Oct. 2019
- ▶ Sequence labeling
  - ▶ [SLP 2<sup>nd</sup> ed.] Ch 5, 6
  - ▶ [SLP 3<sup>rd</sup> ed.] Ch 8, 9
- ▶ Parsing
  - ▶ [AIMA] Ch.23
  - ▶ [SLP 2<sup>nd</sup> ed.] Ch 12, 13, 14
  - ▶ [SLP 3<sup>rd</sup> ed.] Ch 12, 13, 14, 15



# Natural Language Processing

---

- ▶ Get computers to perform useful and interesting tasks involving human languages.
  - ▶ Understanding
  - ▶ Generation
- ▶ Big applications
  - ▶ Question answering, conversational agents (ChatBot)
  - ▶ Financial document processing
  - ▶ Machine translation
  - ▶ News generation



# Levels of NLP Research

---

Phonetics and phonology	knowledge about linguistic sounds
Morphology	knowledge of the meaningful components of words
Syntax	knowledge of the structural relationships between words
Lexical semantics	knowledge of word meaning
Compositional semantics	knowledge of the meaning of sentences
Pragmatics	knowledge of the relationship of meaning to the goals and intentions of the speaker
Discourse	knowledge about linguistic units larger than a single sentence





# Sequence Labeling

# Sequence Labeling

---

- ▶ Problem Definition

- ▶ Known

- ▶ A set of labels  $\mathcal{C} = \{c_1, c_2, \dots, c_J\}$

- ▶ Input

- ▶ Sentence  $s = \{x^1, x^2, \dots, x^m\}$

- ▶ Output

- ▶ For each word  $x^i$ , predict a label  $c^i \in \mathcal{C}$



# Examples

---

## ▶ Part-of-speech tagging

### ▶ Input

Pierre Vinken , 61 years old , will join ...

### ▶ Output

NNP NNP , CD NNS JJ , MD VB

NNP = Proper noun, singular

CD = Cardinal number

NNS = Noun, plural

JJ = Adjective

...



# Examples

---

## ▶ Chinese word segmentation

### ▶ Input

瓦 里 西 斯 的 船 只 中 ...

### ▶ Output

B	I	I	E	S	B	E	S	
(瓦	里	西	斯)	(的)	(船	只)	(中)	...

B = beginning of a word

I = inside of a word

E = end of a word

S = single character word





# Examples

---

## ▶ Named entity recognition

### ▶ Input

Michael Jeffrey Jordan was born in Brooklyn ...

### ▶ Output

B-PER	I-PER	E-PER	O	O	O	S-LOC
<u>Michael Jeffrey Jordan</u>						<u>Brooklyn</u>
Person						Location

B = beginning of an entity

I = inside of an entity

E = end of an entity

S = single word entity

O = outside of any entity

-PER = person

-LOC = location

-ORG = organization

...



# Examples

---

## ▶ Semantic role labeling

### ▶ Input

The cat loves hats ...

### ▶ Output

B-ARG0 E-ARG0 S-PRED S-ARG1  
The cat ← arg0 loves → arg1 hats

B = beginning of an entity

I = inside of an entity

E = end of an entity

S = single word entity

O = outside of any entity

-PRED = predicate

-ARG0 = agent

-ARG1 = patient

...



# The simplest method

---

- ▶ For each word, predict its most frequent label
  - ▶ 90% accuracy on POS tagging!
  - ▶ Disadvantages:
    1. It does not consider the contextual info
      - ▶ “book a flight” vs. “read a book”
      - ▶ 我骑车差点摔倒，好在我一把把把把住了
    2. It does not consider relations between adjacent labels
      - ▶ In BIOES: "B-I" and "B-E" are OK, but "B-O" and "B-S" are not



# Methods

---

- ▶ Hidden Markov Models (HMM)
- ▶ Max-Entropy Markov Models (MEMM)
- ▶ Conditional Random Fields (CRF)



# Hidden Markov Model (HMM)


---

## ▶ Variables

- ▶ X: word
- ▶ Y: label (hidden state)

## ▶ Parameters

- ▶ Transition model  $P(y_t|y_{t-1})$
- ▶ Emission model  $P(x_t|y_t)$
- ▶ Initial distribution  $P(y_1)$ 
  - ▶ Can be seen as transition from  $Y_0=\text{START}$  to  $Y_1$
- ▶ Final distribution  $P(y_n)$ 
  - ▶ Can be seen as transition from  $Y_n$  to  $Y_{n+1}=\text{STOP}$



# HMM Example


---

**Transition**

$Y_{t-1}$	$P(Y_t Y_{t-1})$			
	N	V	P	...
START	0.5	0.1	0.1	...
N	0.4	0.3	0.1	...
V	0.5	0	0.3	...
P	0.3	0.1	0	...
...	...	...	...	...

**Emission**

$Y_t$	$P(X_t Y_t)$			
	"fox"	"dog"	"run"	...
N	0.02	0.03	0.01	...
V	0	0	0.05	...
P	0	0	0	...
...	...	...	...	...



# HMM Inference

---

- ▶ Find the most likely label sequence of the input sentence
  - ▶  $\arg \max_{y_{0:t}} P(y_{0:t} | x_{1:t})$
- ▶ Algorithm?
  - ▶ Viterbi algorithm


$$\begin{aligned} \mathbf{m}_{1:t+1} &= \text{VITERBI}(\mathbf{m}_{1:t}, e_{t+1}) \\ &= P(e_{t+1} | X_{t+1}) \max_{x_t} P(X_{t+1} | x_t) \mathbf{m}_{1:t}[x_t] \end{aligned}$$



# Beyond HMM

---


- ▶ The simplest method: for each word, predict its most frequent label
  - ▶ Problems:
    - ☹️ 1. It does not consider the contextual info
    - 😊 2. It does not consider relations between adjacent labels
- ▶ HMM handles problem 2, but not 1





# Max-Entropy Markov Models (MEMM)

---



$$P(y_{1:n}|x_{1:n}) = P(y_1|x_1) \prod_{t=2}^n P(y_t|y_{t-1}, x_t)$$

$$P(y_t|y_{t-1}, x_t) = \frac{\exp(W^T f(y_{t-1}, y_t, x_t))}{Z(y_{t-1}, x_t)}$$


Possible features:

- $y_{t-1}$  is B and  $y_t$  is E?
- $y_{t-1}$  is B and  $y_t$  is O?
- $x_t$  is a noun?
- $x_t$  is capitalized?
- ...



# Max-Entropy Markov Models (MEMM)

---



$$P(y_{1:n}|x_{1:n}) = P(y_1|x_{1:n}) \prod_{t=2}^n P(y_t|y_{t-1}, x_{1:n})$$


$$P(y_t|y_{t-1}, x_{1:n}) = \frac{\exp(W^T f(y_{t-1}, y_t, x_{1:n}))}{Z(y_{t-1}, x_{1:n})}$$

- ▶ MEMM considers both contextual info and relations between adjacent labels!
- ▶ But... MEMM suffers from **label bias problem**



# Label Bias Problem

---




- ▶ What the local transition probabilities say:
  - ▶ Label 1 prefers to go to label 2
  - ▶ Label 2 prefers to stay at label 2



# Label Bias Problem

---




- ▶  $P(1 \rightarrow 1 \rightarrow 1 \rightarrow 1) = 0.4^3 = 0.064$
- ▶  $P(1 \rightarrow 2 \rightarrow 1 \rightarrow 2) = 0.6 * 0.1 * 0.6 = 0.036$

- ▶  $P(2 \rightarrow 2 \rightarrow 2 \rightarrow 2) = 0.3^3 = 0.027$
- ▶  $P(2 \rightarrow 1 \rightarrow 2 \rightarrow 1) = 0.1 * 0.6 * 0.1 = 0.006$



# Label Bias Problem

---




- ▶ Label 1 has only two transitions but label 2 has five
  - ▶ Transition probabilities from label 2 are lower
- 



# Label Bias Problem

---




- ▶ Label bias in MEMM
  - ▶ Preference of states with lower number of transitions



# Label Bias Problem

---




- ▶ Solution
  - ▶ From local probabilities to local potentials



# From MEMM to CRF

---



$$P(y_{1:n}|x_{1:n}) = P(y_1|x_{1:n}) \prod_{t=2}^n P(y_t|y_{t-1}, x_{1:n})$$


$$P(y_t|y_{t-1}, x_{1:n}) = \frac{\exp(W^T f(y_{t-1}, y_t, x_{1:n}))}{Z(y_{t-1}, x_{1:n})}$$





# From MEMM to CRF

---



$$P(y_{1:n}|x_{1:n}) = \frac{1}{Z(x_{1:n})} \prod_{t=1}^n \exp(W^T f(y_{t-1}, y_t, x_{1:n}))$$

- ▶ Conditional Random Field (CRF) is an undirected graphical model
  - ▶ Global normalization instead of local normalization
  - ▶ Inference: Viterbi



# Summary

---

- ▶ Sequence labeling
  - ▶ Predict a label for each word of a sentence
  - ▶ Many NLP tasks can be seen as sequence labeling
- ▶ Methods
  - ▶ HMM
  - ▶ MEMM
  - ▶ CRF





# Parsing



# Formal Grammars



# Constituency

---

- ▶ Constituents

- ▶ Groups of words within sentences can be shown to act as single units.

- ▶ Ex: (The fox)(jumps (over (the dog)))

- ▶ These units form coherent classes

- ▶ Units in the same class behave in similar ways
    - ▶ ...with respect to their *internal* structure
    - ▶ ...and with respect to other (*external*) units in the language
  - ▶ E.g., noun phrases



# Constituency

---

- ▶ For example, it makes sense to say that the following are all *noun phrases* in English...

Harry the Horse  
the Broadway coppers  
they

a high-class spot such as Mindy's  
the reason he comes into the Hot Box  
three parties from Brooklyn

- ▶ Why?
  - ▶ Similar internal structures
    - ▶ e.g., determiner + modifier + noun + modifier
  - ▶ They can all precede verbs (external evidence)



# Grammars and Constituency

---

- ▶ Grammar
  - ▶ the set of constituents and the rules that govern how they combine
- ▶ Lots of different theories of grammar
- ▶ Context-free grammars (CFGs)
  - ▶ Also known as: Phrase structure grammars
  - ▶ One of the simplest and most basic grammar formalisms



# Context-Free Grammars

---

- ▶ A context-free grammar has four components
  - ▶ A set  $\Sigma$  of terminals (words)
  - ▶ A set  $N$  of nonterminals (phrases)
  - ▶ A start symbol  $S \in N$
  - ▶ A set  $R$  of production rules
    - ▶ Specifies how a nonterminal can produce a string of terminals and/or nonterminals





# Example Grammar

Grammar Rules	Examples
$S \rightarrow NP VP$	I + want a morning flight
$NP \rightarrow$   <i>Pronoun</i>   <i>Proper-Noun</i>   <i>Det Nominal</i> $Nominal \rightarrow$ <i>Nominal Noun</i>   <i>Noun</i>	I Los Angeles a + flight morning + flight flights
$VP \rightarrow$   <i>Verb</i>   <i>Verb NP</i>   <i>Verb NP PP</i>   <i>Verb PP</i>	do want + a flight leave + Boston + in the morning leaving + on Thursday
$PP \rightarrow$ <i>Preposition NP</i>	from + Los Angeles



# Example Grammar

---

*Noun* → *flights* | *breeze* | *trip* | *morning*  
*Verb* → *is* | *prefer* | *like* | *need* | *want* | *fly*  
*Adjective* → *cheapest* | *non-stop* | *first* | *latest*  
                  | *other* | *direct*  
*Pronoun* → *me* | *I* | *you* | *it*  
*Proper-Noun* → *Alaska* | *Baltimore* | *Los Angeles*  
                  | *Chicago* | *United* | *American*  
*Determiner* → *the* | *a* | *an* | *this* | *these* | *that*  
*Preposition* → *from* | *to* | *on* | *near*  
*Conjunction* → *and* | *or* | *but*



# Sentence Generation

---

- ▶ A grammar can be used to generate a string
  - ▶ starting from a string containing only the start symbol  $S$
  - ▶ recursively applying the rules to rewrite the string
  - ▶ until the string contains only terminals
- ▶ The generative process specifies the **grammatical structure (parse tree)** of the string



# Example

$S \rightarrow NP VP$

$S \rightarrow Aux NP VP$

$S \rightarrow VP$

$NP \rightarrow Pronoun$

$NP \rightarrow Proper-Noun$

$NP \rightarrow Det Nominal$

$NP \rightarrow Nominal$

$Nominal \rightarrow Noun$

$Nominal \rightarrow Nominal Noun$

$Nominal \rightarrow Nominal PP$

$VP \rightarrow Verb$

$VP \rightarrow Verb NP$

$VP \rightarrow Verb NP PP$


$VP \rightarrow Verb PP$

$VP \rightarrow Verb NP NP$

$VP \rightarrow VP PP$

$PP \rightarrow Preposition NP$

.....



*Book the dinner flight*



# Sentence Parsing

---

- ▶ Parsing is the process of taking a string and a grammar and returning one or more parse tree(s) for that string
  - ▶ If no parse tree can be found, then the string does not belong to the language
  - ▶ Parsing algorithms: CYK, Earley, etc.
    - ▶ To be introduced later



# Probabilistic Grammars

---

- ▶ Also called stochastic grammars
- ▶ Each rule is associated with a probability

$$\alpha \rightarrow \beta : P(\alpha \rightarrow \beta | \alpha)$$


- ▶ The probability of a parse tree is the product of the probabilities of all the rules used in generating the parse tree



# Example

$S \rightarrow NP VP$	[.80]
$S \rightarrow Aux NP VP$	[.15]
$S \rightarrow VP$	[.05]
$NP \rightarrow Pronoun$	[.35]
$NP \rightarrow Proper-Noun$	[.30]
$NP \rightarrow Det Nominal$	[.20]
$NP \rightarrow Nominal$	[.15]
$Nominal \rightarrow Noun$	[.75]
$Nominal \rightarrow Nominal Noun$	[.20]
$Nominal \rightarrow Nominal PP$	[.05]
$VP \rightarrow Verb$	[.35]
$VP \rightarrow Verb NP$	[.20]
$VP \rightarrow Verb NP PP$	[.10]
$VP \rightarrow Verb PP$	[.15]
$VP \rightarrow Verb NP NP$	[.05]
$VP \rightarrow VP PP$	[.15]
$PP \rightarrow Preposition NP$	[1.0]

.....



*Book the dinner flight*

$$P(T) = .05 \times .20 \times .20 \times .20 \times .75 \times .30 \times .60 \times .10 \times .40 = 2.2 \times 10^{-6}$$



# Ambiguity

---

- ▶ A sentence is ambiguous if it has more than one possible parse tree
  - ▶ ...and hence more than one interpretation
- ▶ Examples
  - ▶ Time flies like an arrow.
  - ▶ Astronomers saw stars with ears.





# Example

---

S → NP VP    1.0

PP → P NP    1.0

VP → V NP    0.7

VP → VP PP    0.3

P → *with*    1.0

V → *saw*    1.0

NP → NP PP    0.4

NP → *astronomers*    0.1

NP → *ears*    0.18

NP → *saw*    0.04



NP → *stars*    0.18

NP → *telescopes*    0.1



# Example

---




$$\begin{aligned} P(t_1) &= 1.0 \times 0.1 \times 0.7 \times 1.0 \times 0.4 \\ &\quad \times 0.18 \times 1.0 \times 1.0 \times 0.18 \\ &= 0.0009072 \end{aligned}$$



# Example

---



$$\begin{aligned} P(t_2) &= 1.0 \times 0.1 \times 0.3 \times 0.7 \times 1.0 \\ &\quad \times 0.18 \times 1.0 \times 1.0 \times 0.18 \\ &= 0.0006804 \end{aligned}$$




# Chomsky Normal Form (CNF)

---

- ▶ Only two types of production rules in CNF

$$A \rightarrow B C$$

$$A \rightarrow w$$



Noam Chomsky

- ▶ Any arbitrary CFG can be rewritten into CNF automatically
  - ▶ The resulting grammar accepts (and rejects) the same set of strings as the original grammar
  - ▶ But the resulting parse trees are different (i.e., binarized)

